


# 2IMN20 - Real-Time Systems

## Scheduling 101

**Geoffrey Nelissen**

2023-2024

# Buttazzo's book, chapters 2 and small part of 3

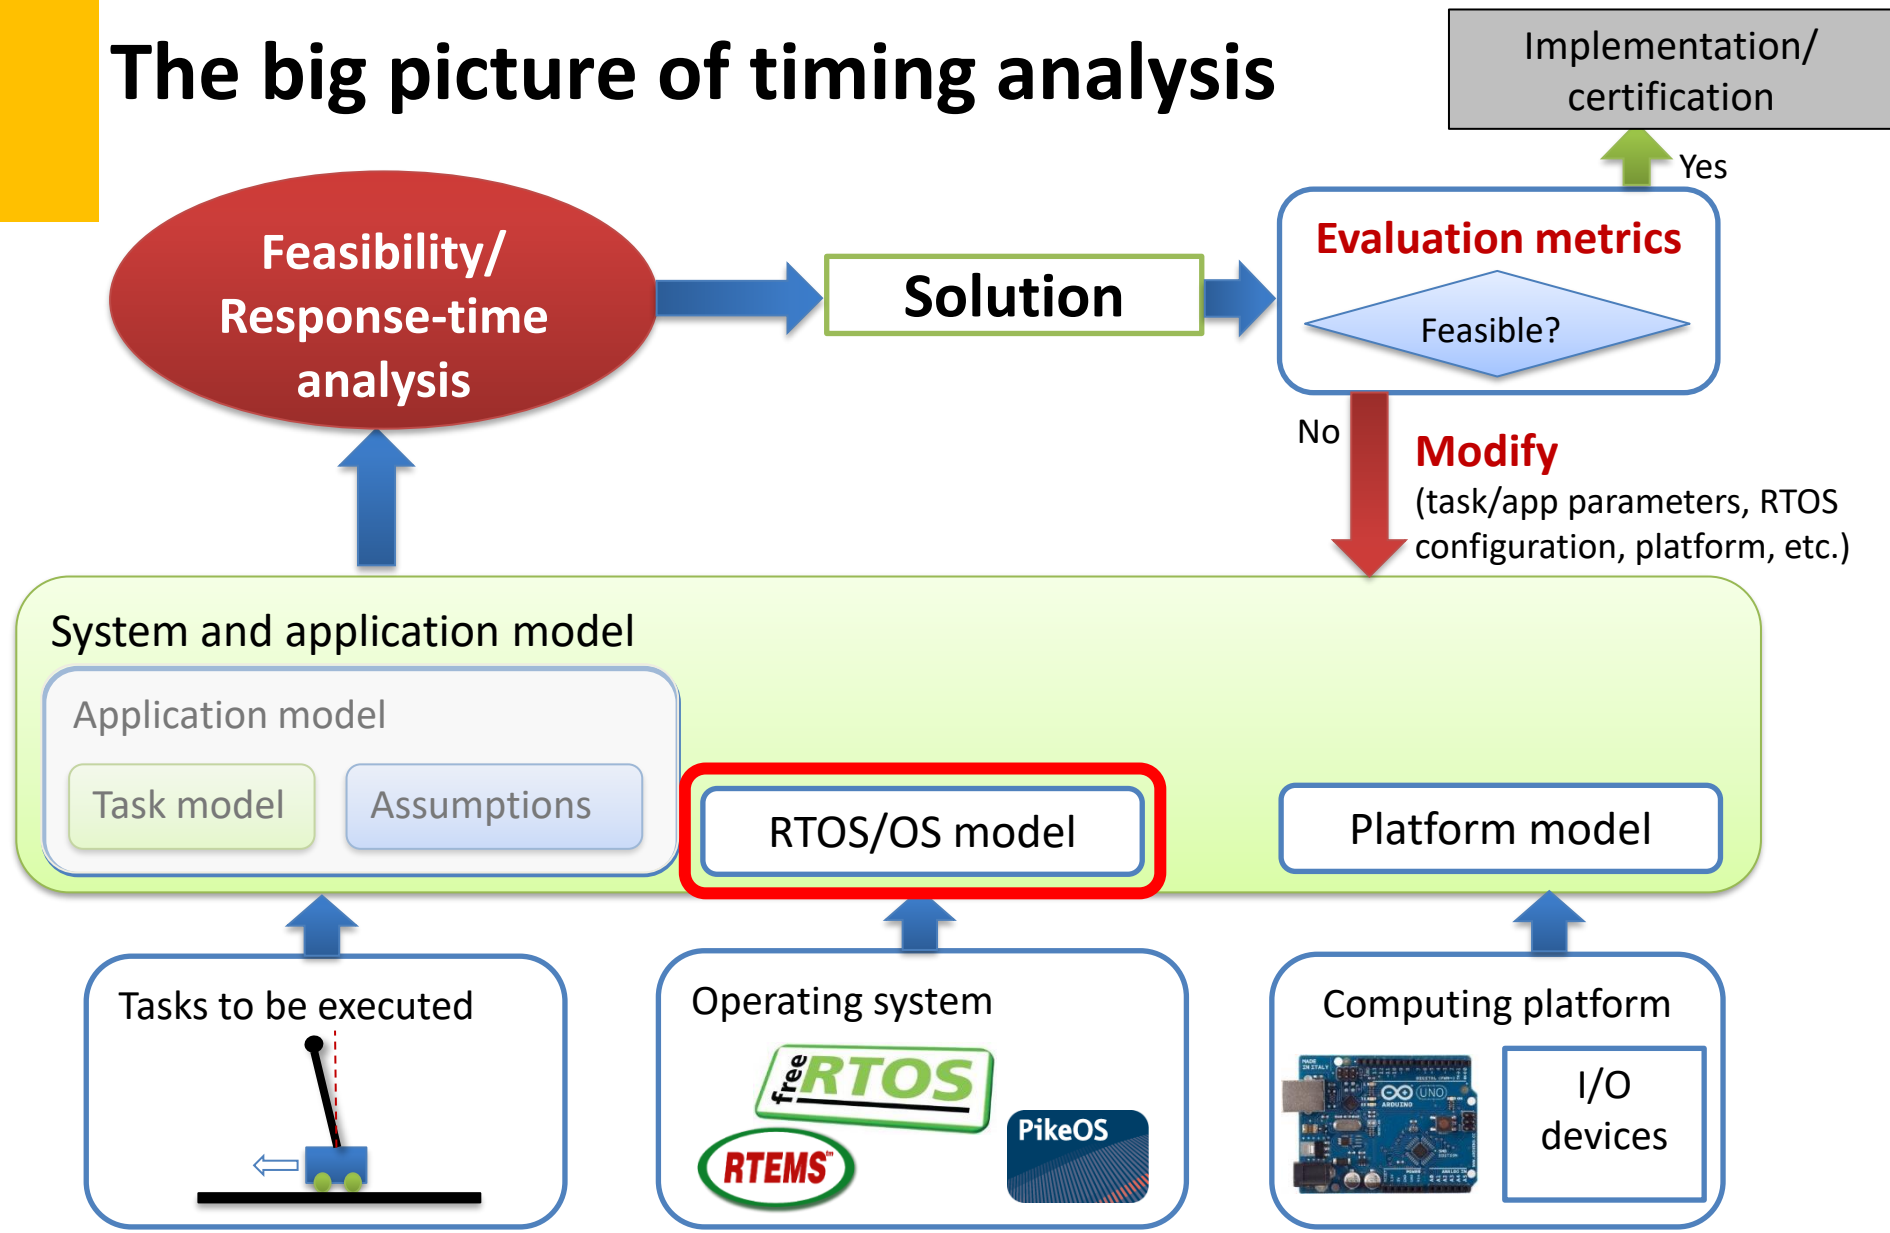


## Disclaimer:

**Most slides were provided by Dr. Mitra Nasri**

Some slides have been taken from [Giorgio Buttazzo](#)

# The big picture of timing analysis

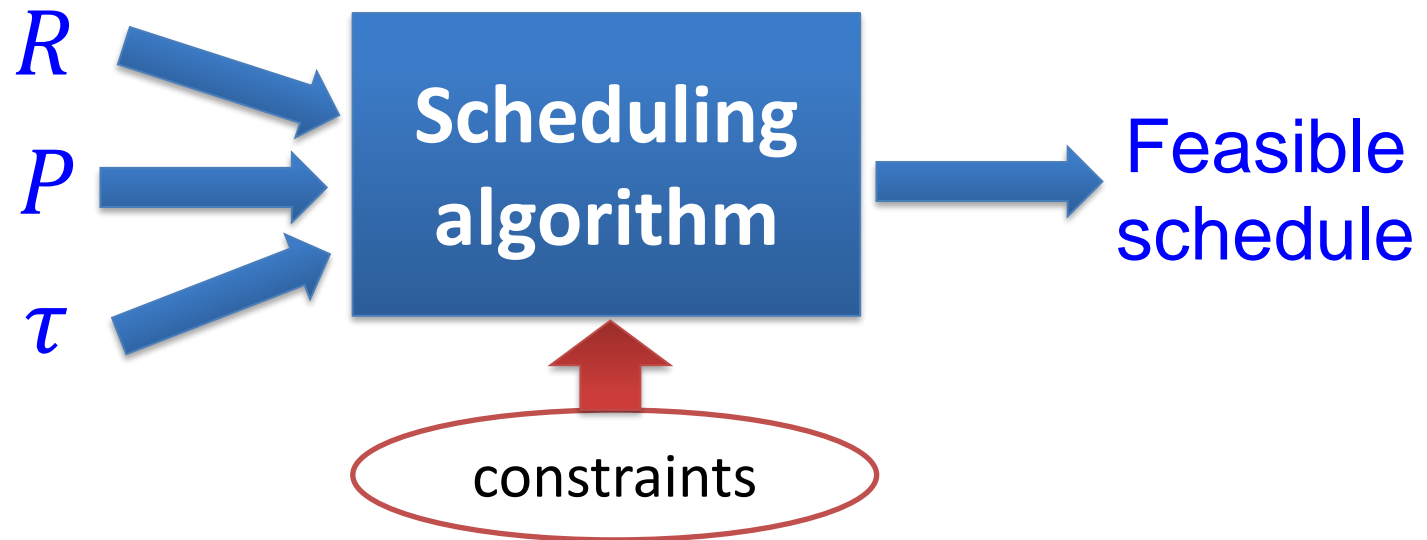


# Agenda

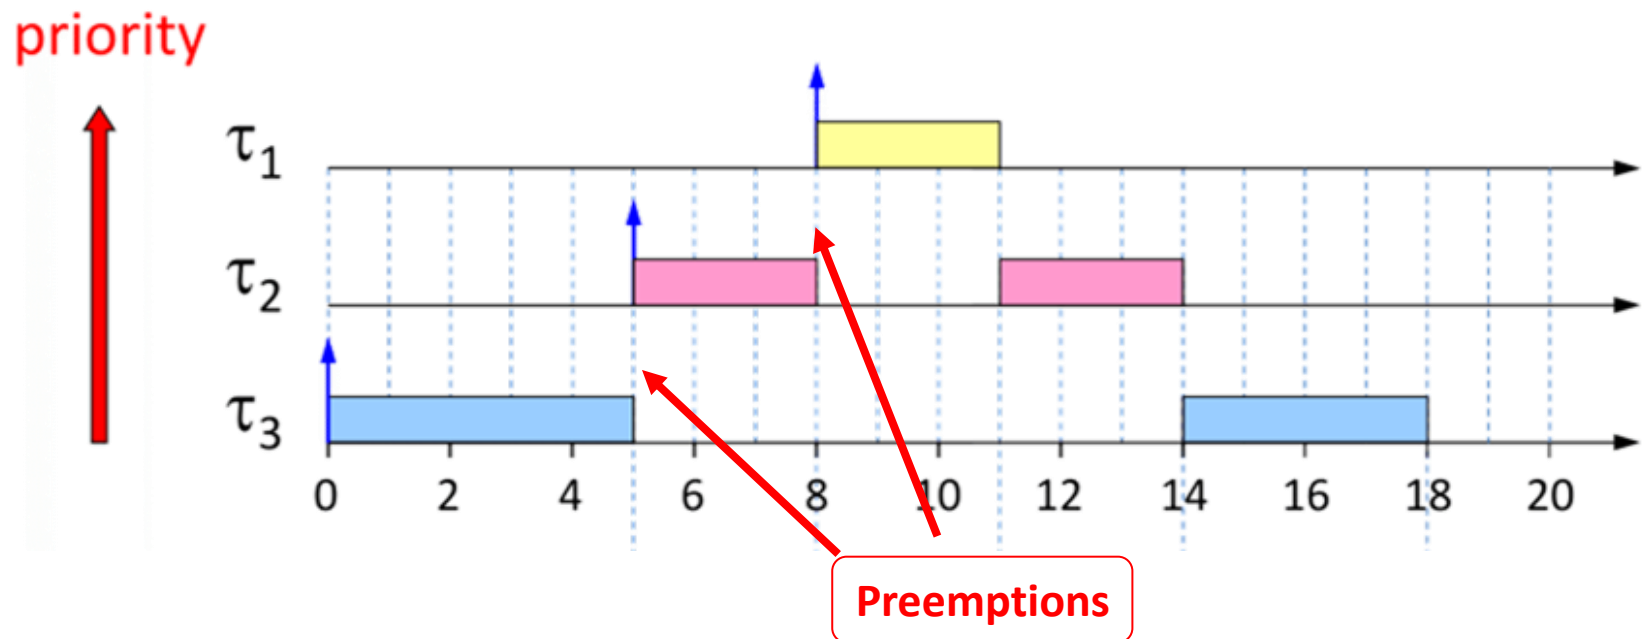
- Scheduling
- Table-driven scheduling
- Online scheduling policies

# General scheduling problem

**Given** a set  $\tau$  of  $n$  tasks, a set  $P$  of  $m$  processors, and a set  $R$  of  $r$  resources,  
**find** an assignment of  $P$  and  $R$  to  $\tau$  for any time instant  $t$  that produces a feasible schedule under a set of constraints.



# Example schedule: FP Preemptive scheduling



# Scheduling policy/algorithm (for tasks on processors)

A **scheduling policy** is an algorithm that determines what task executes when on what processor, i.e.,  
**produces a schedule.**

# Definitions

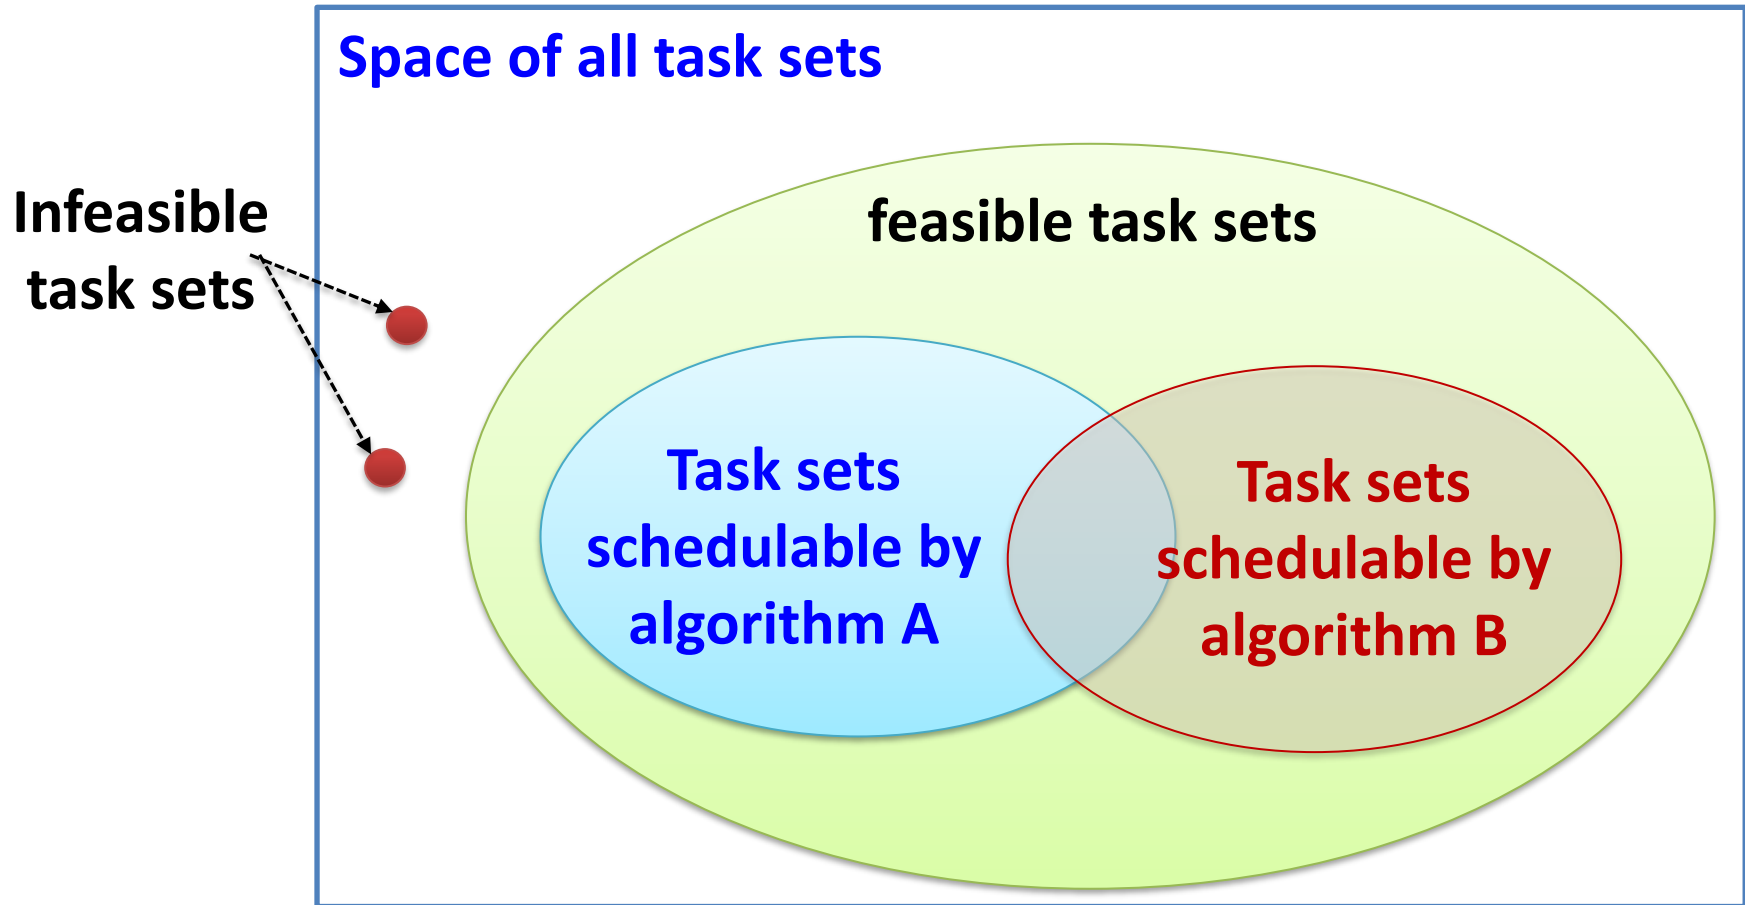
A schedule  $S$  is said to be **feasible** if it satisfies all given requirements, for example, “no deadline is missed in  $S$ ”.

A task set  $\tau$  is said to be **feasible**, if there always exists at least one feasible schedule for it

A task set  $\tau$  is said to be **schedulable** with a **scheduling algorithm**  $A$ , if  $A$  always generates a feasible schedule for  $\tau$ .



# Feasibility vs. schedulability



# Properties of scheduling algorithms

- **Preemptive vs. non-preemptive**
- **Work-conserving vs. non-work-conserving**
- **Offline vs. online**
- **Optimal vs. non-optimal**

# Properties of scheduling algorithms

- Preemptive vs. non-preemptive
- **Work-conserving** vs. **non-work-conserving**
- **Offline** vs. **online**
- **Optimal** vs. **non-optimal**

# Work-conserving vs. non-work-conserving

- **Work-conserving**

Such algorithm **does not** leave the processor **idle** as long as there is a ready task in the system (a task is in the ready queue).

- **Non-work conserving**


Such algorithm **may** leave the processor **idle** even if there is a ready task in the ready queue.

**Q:** Why would we want to leave the processor idle when there is something to execute?

See next slide

# Work-conserving vs. non-work-conserving

Q: Is this job set feasible under non-preemptive scheduling?



The processor must remain idle in  $[3, 5)$  even though Tasks 2 and 3 are in the ready queue

→ We must use a non-work-conserving algorithm

# Optimal v.s. non-optimal

- **Optimal**

They generate a schedule that minimizes a [cost function](#), defined based on an [optimality criterion](#).

- **Non-optimal (heuristic)**

They generate a schedule according to a heuristic function that tries to satisfy an optimality criterion, but [there is no guarantee of success](#).

## Example for some optimization goals

- **Feasibility**: Find a feasible schedule if there exists one.
- Minimize the **maximum lateness (i.e., tardiness)**
- Minimize the **number of deadline misses**
- Minimize the **average waiting time**
- Assign a value to each task, then maximize the **cumulative value** of the feasible tasks

# Offline vs. online

- **Offline**

all scheduling decisions are taken before tasks activations: the schedule is stored in a table (**table-driven scheduling**).

- **Online**

scheduling decisions are taken at run time based on the set of active tasks and the system state.

# **A closer look at offline scheduling policies**



# Table-driven scheduling

- Store the schedule in a table that is prepared offline
- Dispatch jobs according to the table

Task	Start time
1	0
2	10
3	12
1	20
4	28

Cycle = 40

# Table-driven scheduling

- Store the schedule in a table that is prepared offline
- Dispatch jobs according to the table

Task	Start time
1	0
2	10
3	12
1	20
4	28




Cycle = 40

Timer interrupt handler:

1. Dispatcher reads the table entry at the current index  $i$  and dispatches the task
2. Increments  $i$
3. Sets the timer interrupt to the “start time” of the current table entry at index  $i$
4. If the last entry is reached. Reset the time at the end of the cycle time, and set index  $i$  to 0

Resulting  
schedule



# Advantages of table-driven scheduling

Any guess?

- **Extremely predictable**
  - Provides full knowledge about when the tasks will be executing
  - Hence, you can easily analyze system performances
- **Easy to certify**
  - Avionics industry uses partly table-driven scheduling
- **Extremely flexible for schedule optimizations**
  - It is easy to include optimization criteria while building the schedule
  - Can handles various system constraints such as precedence constraints, etc.
- **Low runtime overhead**
  - A true  $O(1)$  algorithm for scheduling on most hardware platforms
- **Very small “code” footprint**
  - Only requires a few instructions to implement

# Disadvantages of table-driven scheduling

Any guess?

- **Requires concrete knowledge of task release times**
  - Cannot be applied on event-based systems or dynamic workloads
- **Building an optimal schedule might be computationally expensive**
  - recall: the general scheduling problem is NP-Hard

**What else?**



# It eats up a large amount of memory to store the table!



## Example:

For a system with 1000 jobs per hyperperiod,  
and 32 bits to store a table entry, the table becomes as big as **4kB**

An **Arduino Mega** has only **8kB of RAM**

# Memory is money!

- Many embedded systems have a limited **processing power** and **memory** because
  - memory is **expensive**
  - consumes energy**



## Arm Cortex MCU family

STM32 32-bit ARM Cortex MCUs	<	STM32F2 Series	STM32F3 Series	STM32F4 Series	STM32F7 Series	S
Total Parts: (752) for STM32 32-bit ARM Cortex MCUs   Matching Parts : (90)						
Part Number	Package ▼	Core	Operating Frequency (MHz) (Processor speed)	FLASH Size (kB) (Prog)	Internal RAM Size (kB)	I/Os (High Current)
STM32L011G4	UFQFPN 28 4x4 x0.55	ARM Co rtex-M...	32	16	2	24
STM32L011K4	LQFP 32 7x7x1.4, ...	ARM Co rtex-M...	32	16	2	28
STM32L021D4	TSSOP 14	ARM Co rtex-M...	32	16	2	11
STM32L021F4	UFQFPN 20 3x3 x0.6	ARM Co rtex-M...	32	16	2	16
STM32L021G4	UFQFPN 28 4x4 x0.55	ARM Co rtex-M...	32	16	2	24
STM32L021K4	LQFP 32 7x7x1.4	ARM Co rtex-M...	32	16	2	28
STM32L031F4	TSSOP 20	ARM Co rtex-M...	32	16	8	15
STM32L071C8	LQFP 48 7x7x1.4	ARM Co rtex-M...	32	64	20	37
STM32L071RZ	LQFP 64 10x10 x1.4, ...	ARM Co rtex-M...	32	192	20	51
STM32L071VB	LQFP 100 14x14 x1.4	ARM Co rtex-M...	32	128	20	84

# Online scheduling policies


## Cyclic executive

See also Buttazzo Section 4.2.

# Single-rate AFAP (As Fast As Possible)

```
while(1){  
    Task_1();  
    Task_2();  
    ...  
    Task_n();  
}
```

A single (infinite) loop, calling each task one after another



## Advantages:

- Simple
- no preemption
- Fast to implement/deploy

## Disadvantages:

- **Timings** (start and finish time of each job) **depends on computation times of *all* tasks;**
  - Potentially *unbounded jitter*, i.e. *drift*
- **Energy inefficient** (always execute even when not needed)
- **No option for “background” tasks**




# Single-rate Time-driven AFAP

```
int k; /* activation counter */

k = -1;
while(1){
    k = k + 1;
    /* wait till absolute time phi + k * T */
    sleepUntil(phi+k*T)
    Task_1();
    Task_2();
    /* ... */
    Task_n();
}
```

Adds a timer to trigger the loop execution periodically



## Advantages:

- Same as Single rate – AFAP
- **Resolves** problems with *drift* and *energy inefficiency* and *background tasks*

## Disadvantage:

- All tasks are all executed at the **same rate**
- We must have  $\sum_{\tau_i \in \tau} C_i \leq T$


# Multi rate Periodic

Division of time in a major and minor cycles

```
int k; /* cycle counter */

k = -1;
while(1){
    k = k + 1;
    /* wait till absolute time phi + k * T_1 */
    sleep(phi + k*T_1);
    Task_1();
    if( k % 2 == 0)
    { Task_2();
    } else {
        Task_3();
    }
}
```

Executes different set of tasks depending on the minor cycle number



## Advantages:

- Same as single-rate time-driven AFAP
- **Reduced** start/finish-jitter of tasks
- **Not** all tasks execute at the **same frequency**
- No need for  $\sum_{\tau_i \in \tau} C_i \leq T$

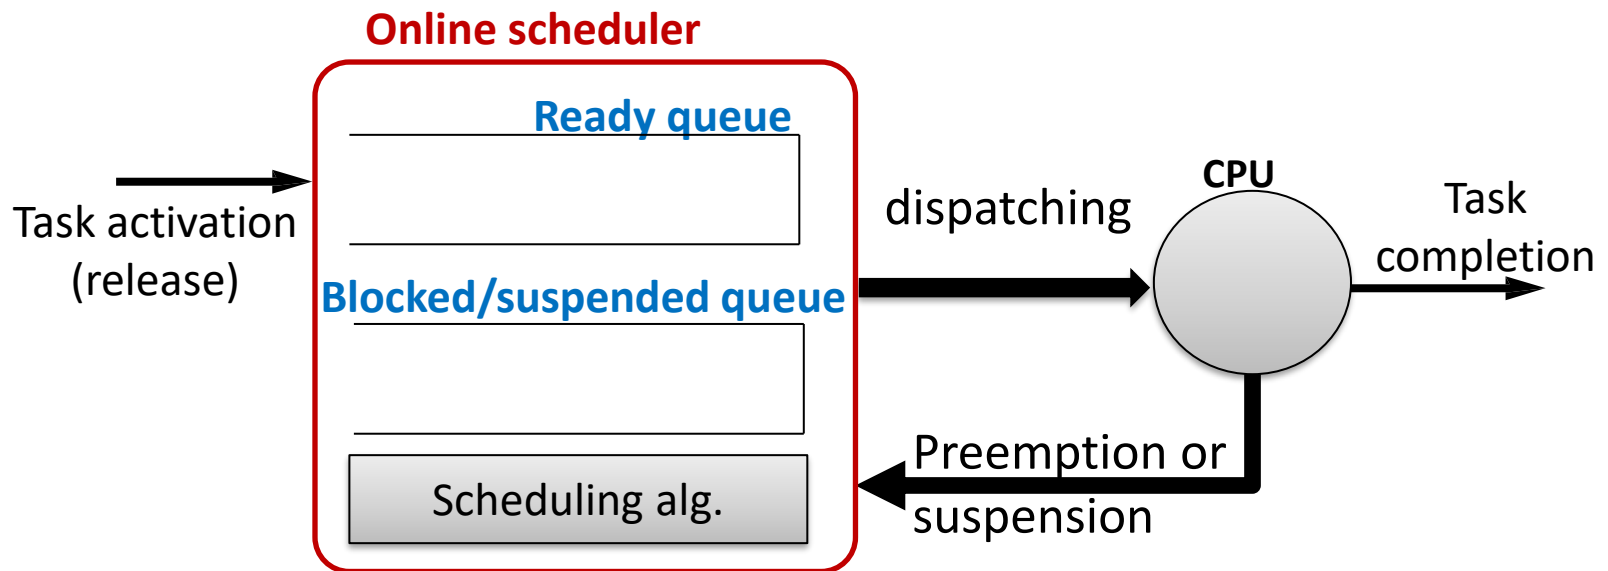
# Summary

- Cyclic executives
  - **Advantages**
    - Fast to implement/deploy, no “preemption costs”, simple:
      - requires just a hardware timer;
      - no need for shared resource access protocols.
  - **Disadvantages**
    - Limited to **periodic** events or **polling** sporadic events
    - Does **not support dynamic** systems where tasks may join and leave during the runtime
    - **Not robust** against overload
    - **Hard to maintain** when (see [Buttazzo], Section 4.2)
      - the frequency of a task must be updated
      - a task added
      - the execution time of a task increases
  - **Examples of use**
    - ROS/ROS2 (the Robot Operating System)
    - Lupo EL (developed by ME in 2009)
    - Signify LED drivers

# **A closer look at other known online scheduling policies**

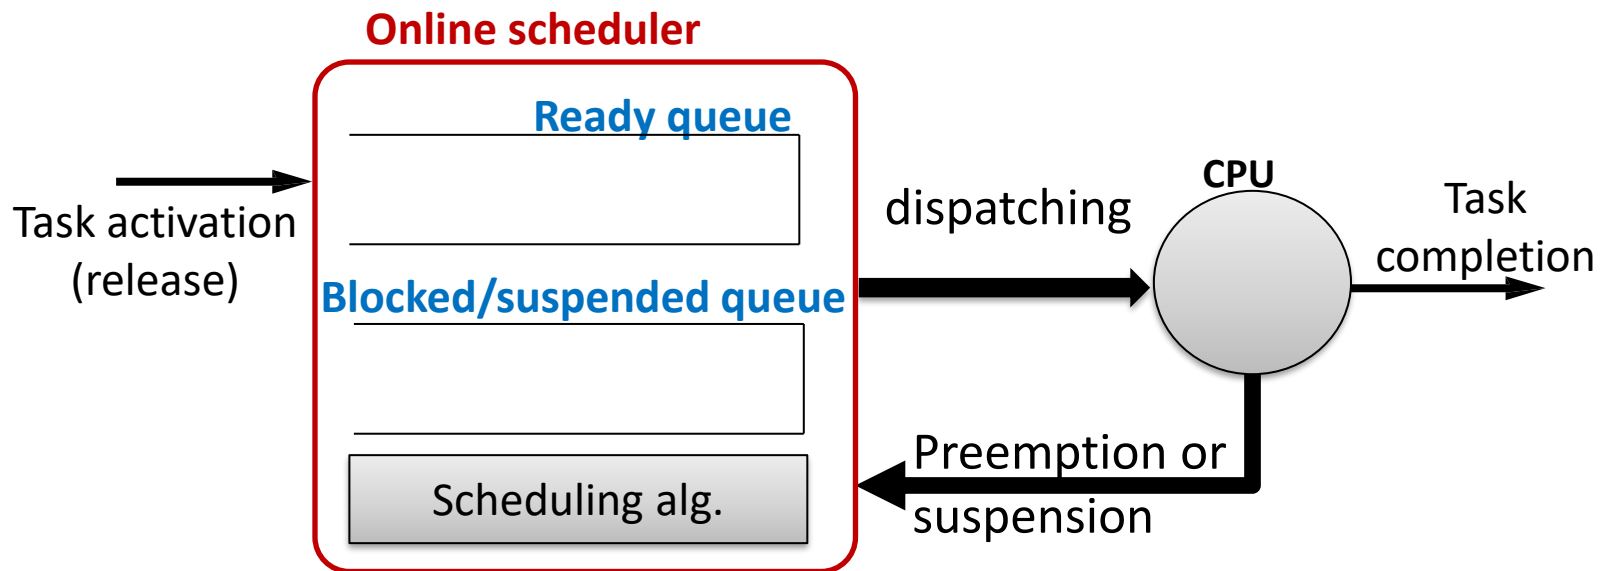
# Online scheduler in an operating system

- In a concurrent system with one processor, several tasks can be simultaneously active, but only one can be in execution (**running**).
- An active task that is **not** in **execution**, **blocked** or **suspended** is said to be **ready**.

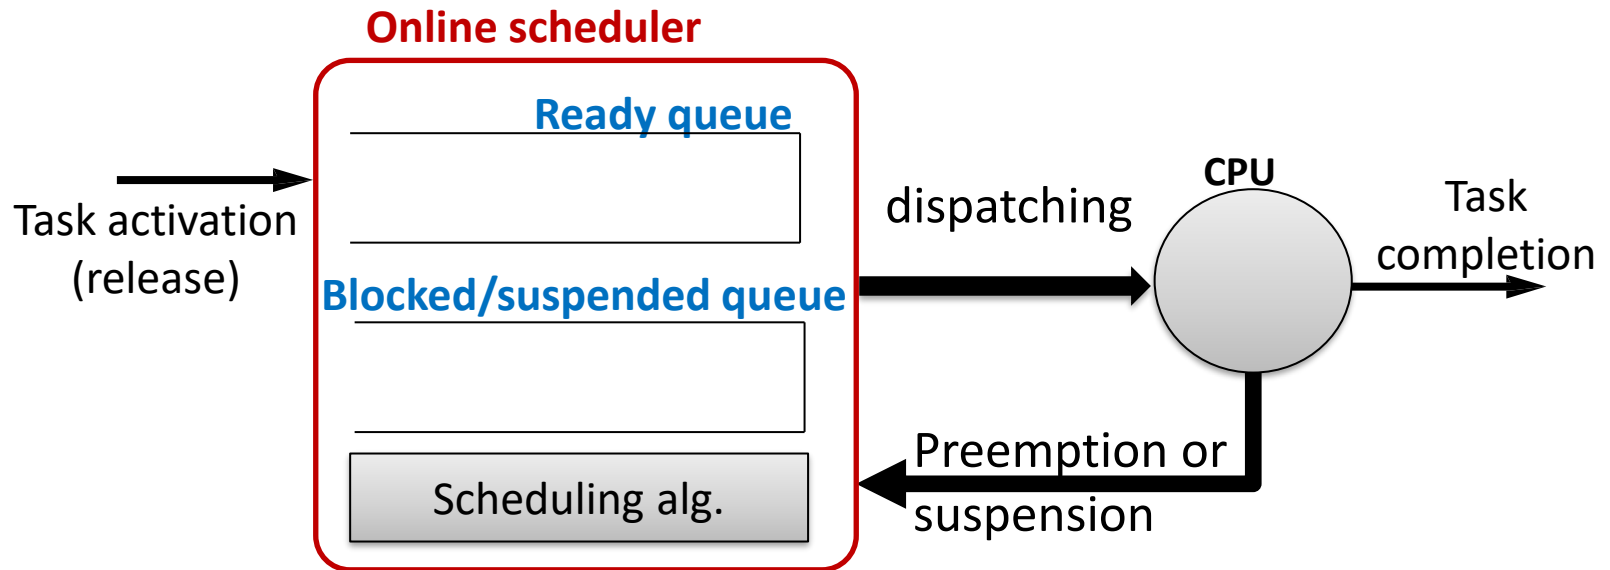


# Online scheduler in an operating system

- Ready tasks are kept in a **ready queue**, managed by a **scheduling policy**.
- The **scheduling policy** determines which task is dispatched on the processor



# Preemption

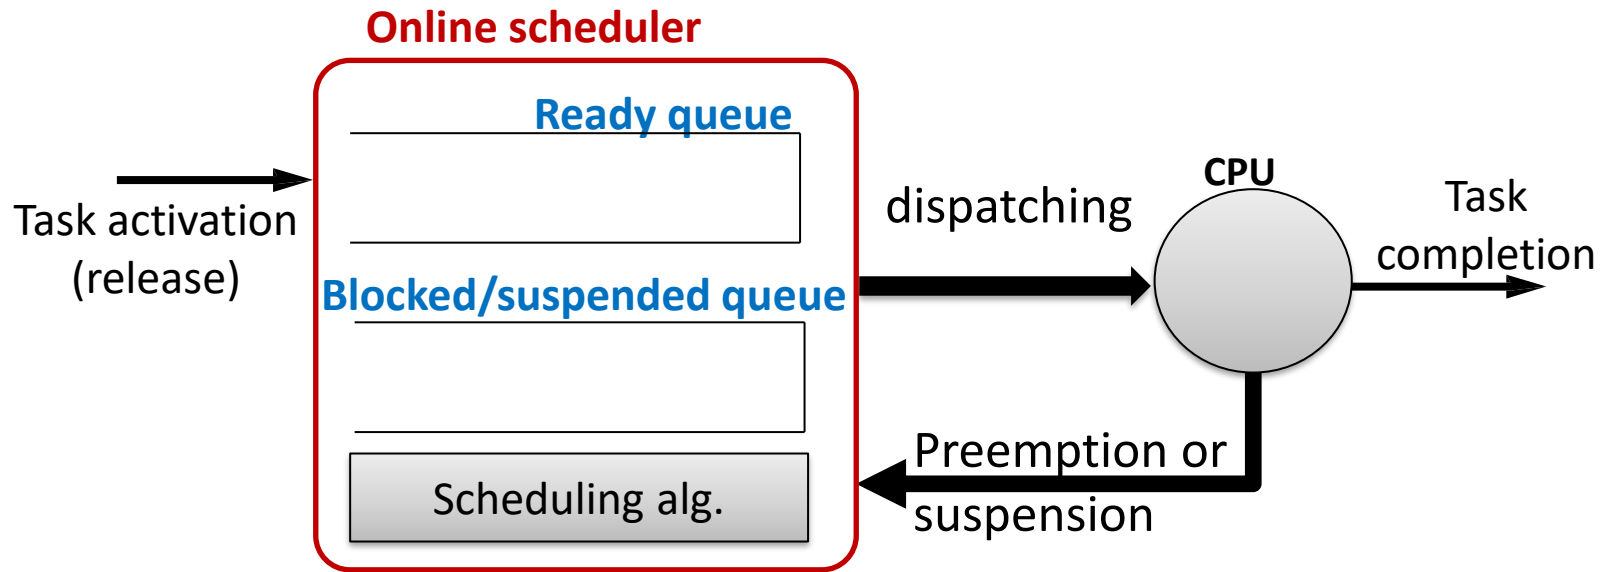


Preemption is a kernel mechanism that allows to preempt the execution of the running task in favor of another task.

The **preempted task** goes back into the **ready queue**.

- Preemption enhances concurrency and may help reducing the response time of high priority tasks.
- It can be disabled (completely or temporarily) to ensure the consistency of certain critical operations.

# Suspension

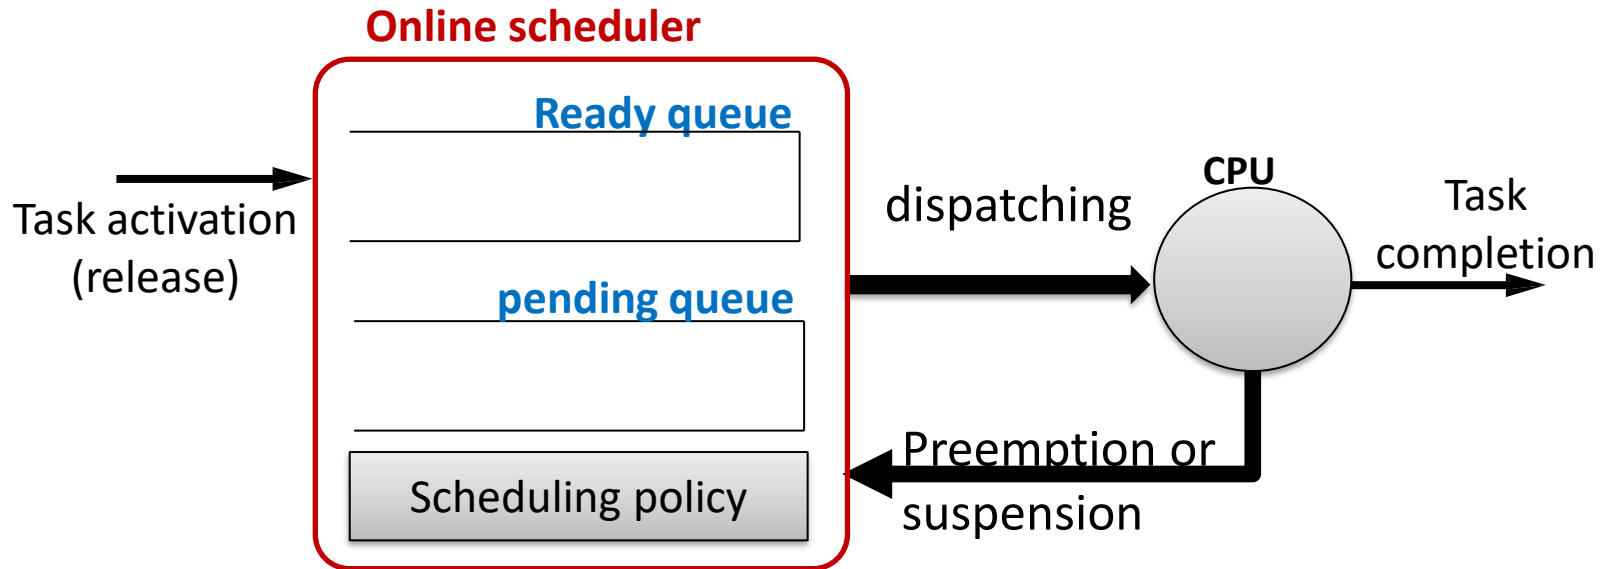


Suspension happens when a task decides to suspend itself (e.g., by “delay( )” or “sleep( )”), or the task makes a non-blocking call to an I/O or message queue or waits for results from another tasks or co-processor.

The **suspended task** goes in the **suspended queue**.



# Online scheduling



**What online scheduling policies (or algorithms) do you know?**

# Some well-known scheduling algorithms

- **First-in-first-out scheduling (FIFO)**  
= First-come-first-serve (FCFS)
- **Round robin**
- **Shortest-job first**
- **Earliest deadline first (EDF)**
- **Fixed-priority scheduling (FP)**

# Online scheduling policies:

## Static vs. dynamic priorities

- Task-level static priorities

- Scheduling decisions are taken based on task's fixed parameters that are known beforehand.
- Examples: Rate Monotonic (RM), Deadline Monotonic (DM)

- Job-level static priorities

- Scheduling decisions are taken based on parameters of the job known only at its release time.
- Examples: FIFO or earliest-deadline first (EDF), which uses the absolute deadline of the jobs in order to decide which job has the highest-priority.


**Note:** a job-level static scheduling policy such as EDF or FIFO is a task-level dynamic priority scheduling policy.

- Job-level dynamic priorities

- Scheduling decisions are taken based on parameters that can change with time.
- Example: scheduling policy with ageing, least-laxity first or shortest remaining execution time first policy

# Example: FIFO (non-preemptive)

- The job with the earliest released is scheduled first




$\tau_i$	$C_i$	$r_{i,1}$	$d_{i,1}$
$\tau_1$	8	5	14
$\tau_2$	3	10	14
$\tau_3$	4	2	17

Assume that each task releases a single job

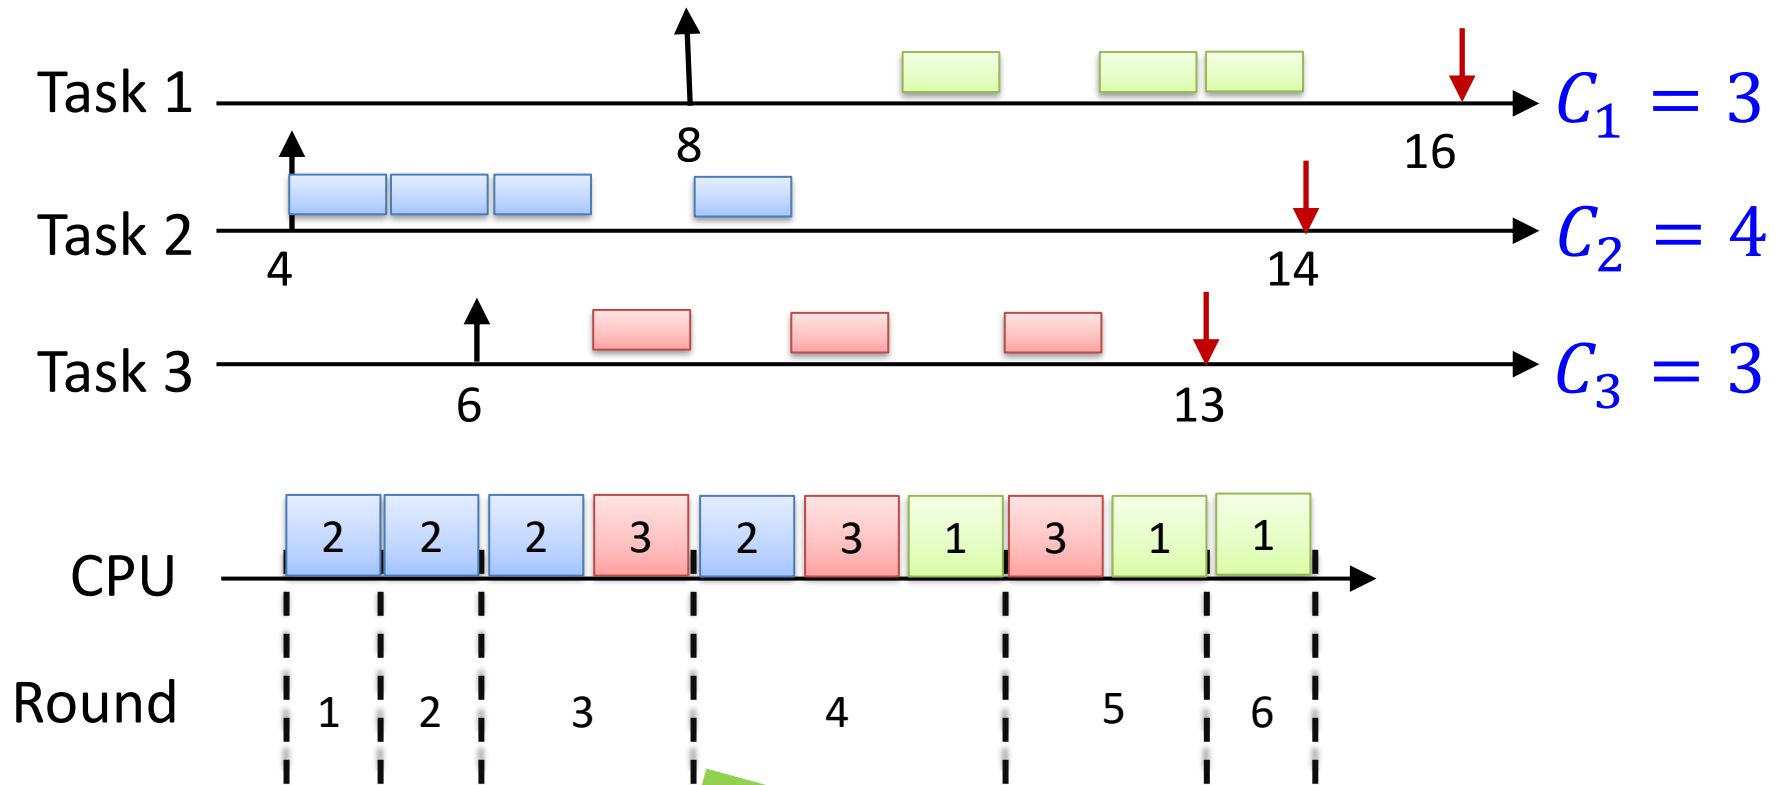
# Round robin

- The ready queue is served with FIFO, but ...
- Each task  $\tau_i$  cannot execute for more than  $Q$  time units ( $Q$  = time quantum).
- When  $Q$  expires,  $\tau_i$  is put back in the queue.




# Round robin

- $n$  = number of task in the ready queue
- Round robin creates “rounds” that are as long as  $n \cdot Q$
- Assume  $Q = 1$



RR generates large number of preemptions

# Example: Shortest Job First (non-preemptive)



$\tau_i$	$C_i$	$r_{i,1}$	$d_{i,1}$
$\tau_1$	8	2	14
$\tau_2$	3	2	14
$\tau_3$	4	2	17


SJF is **difficult to implement** in practice because it **requires to estimate the execution time** of each job, which may be different from their worst-case execution times

Assume that each task releases a single job

# Fixed-priority scheduling (preemptive)

- Each task has a priority  $P_i$ , typically  $P_i \in [0, 255]$   
( $P_i > P_j$  means that Task  $\tau_i$  has a higher priority than task  $\tau_j$ )
- The task with the highest priority is selected for execution.
- Tasks with the same priority are served in FIFO order

$$P_1 > P_2 > P_3$$






# Earliest deadline first (EDF)

- **Algorithm** [Horn 74]
  - Order the ready queue by **increasing absolute deadlines** (job-level fixed priority).

# Example of schedule with EDF



# Earliest deadline first (EDF)

- **Algorithm** [Horn 74]

- Order the ready queue by **increasing absolute deadline (job-level fixed priority)**.

- **Assumptions**

- Horn's algorithm is **preemptive** and is for **independent tasks** executed on a **single core** platform

- **Property**

- Under the assumptions above, EDF **minimizes the maximum lateness ( $L_{max}$ )**

$$L_{max} = \max\{L_{i,j} \mid \forall J_{i,j} \in \text{job set}\}$$

→ EDF is optimal from a feasibility viewpoint

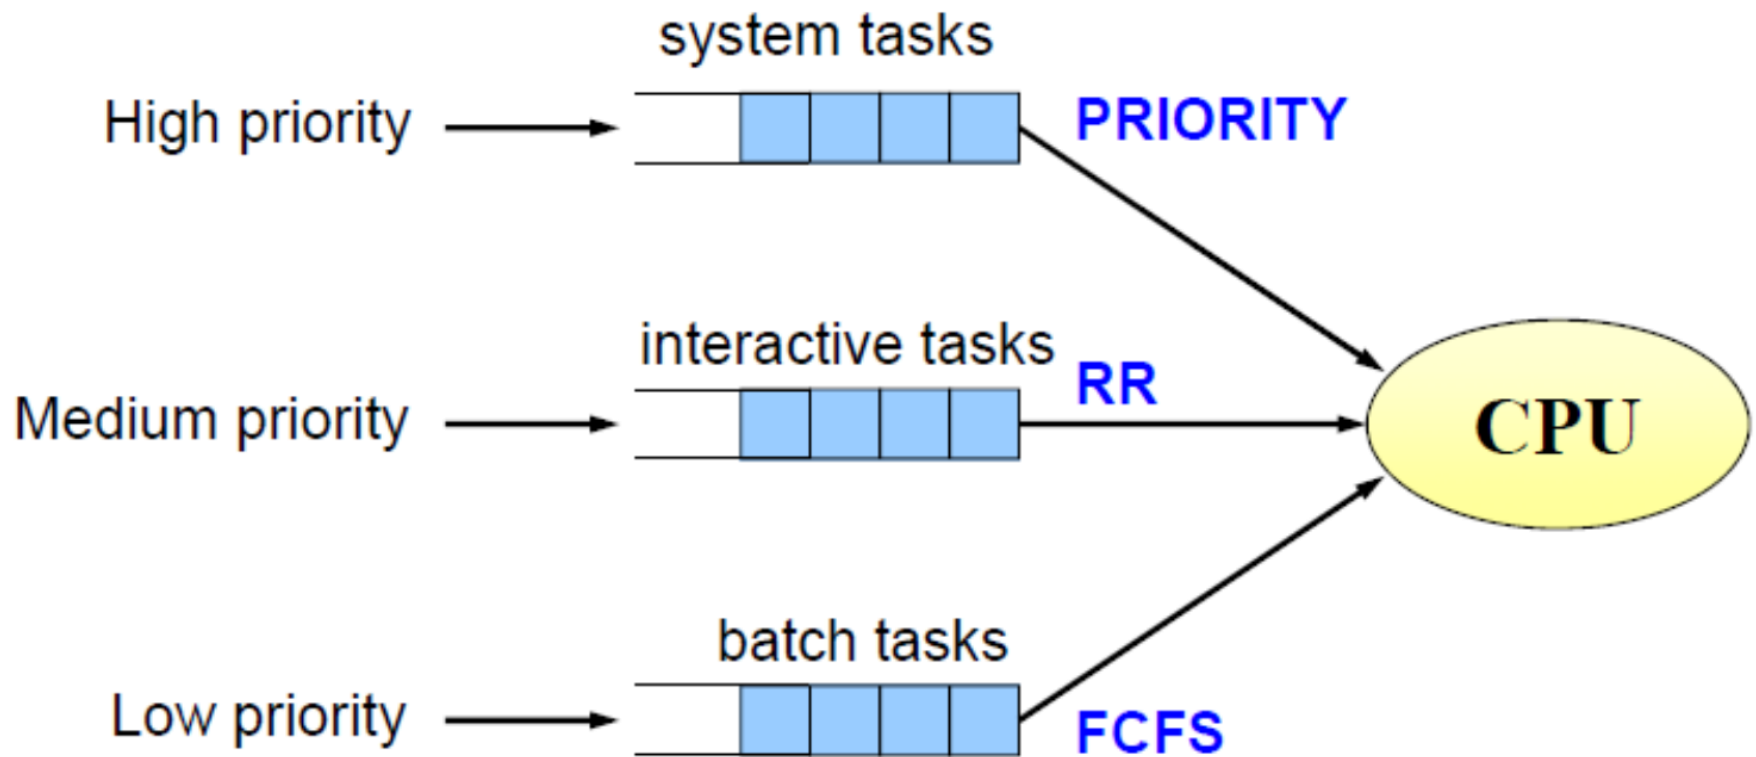
# A property of optimal algorithms

If a task set is not schedulable by an **optimal algorithm**, then it **cannot** be schedulable by **any other algorithm**.

If an algorithm A **minimizes  $L_{max}$**  then A is also **optimal** in the sense of feasibility.  
The opposite is not true.

➔ EDF is therefore optimal (w.r.t. feasibility) for the scheduling independent tasks on single core

# Multi-level scheduling



# Summary

- FIFO is fast and simple to implement, but bad at guaranteeing deadlines
- Round robin generates a large number of preemptions
- EDF is optimal w.r.t. feasibility (because it minimizes the maximum lateness)
- Table-driven scheduling is inflexible and trades computing complexity for memory complexity