

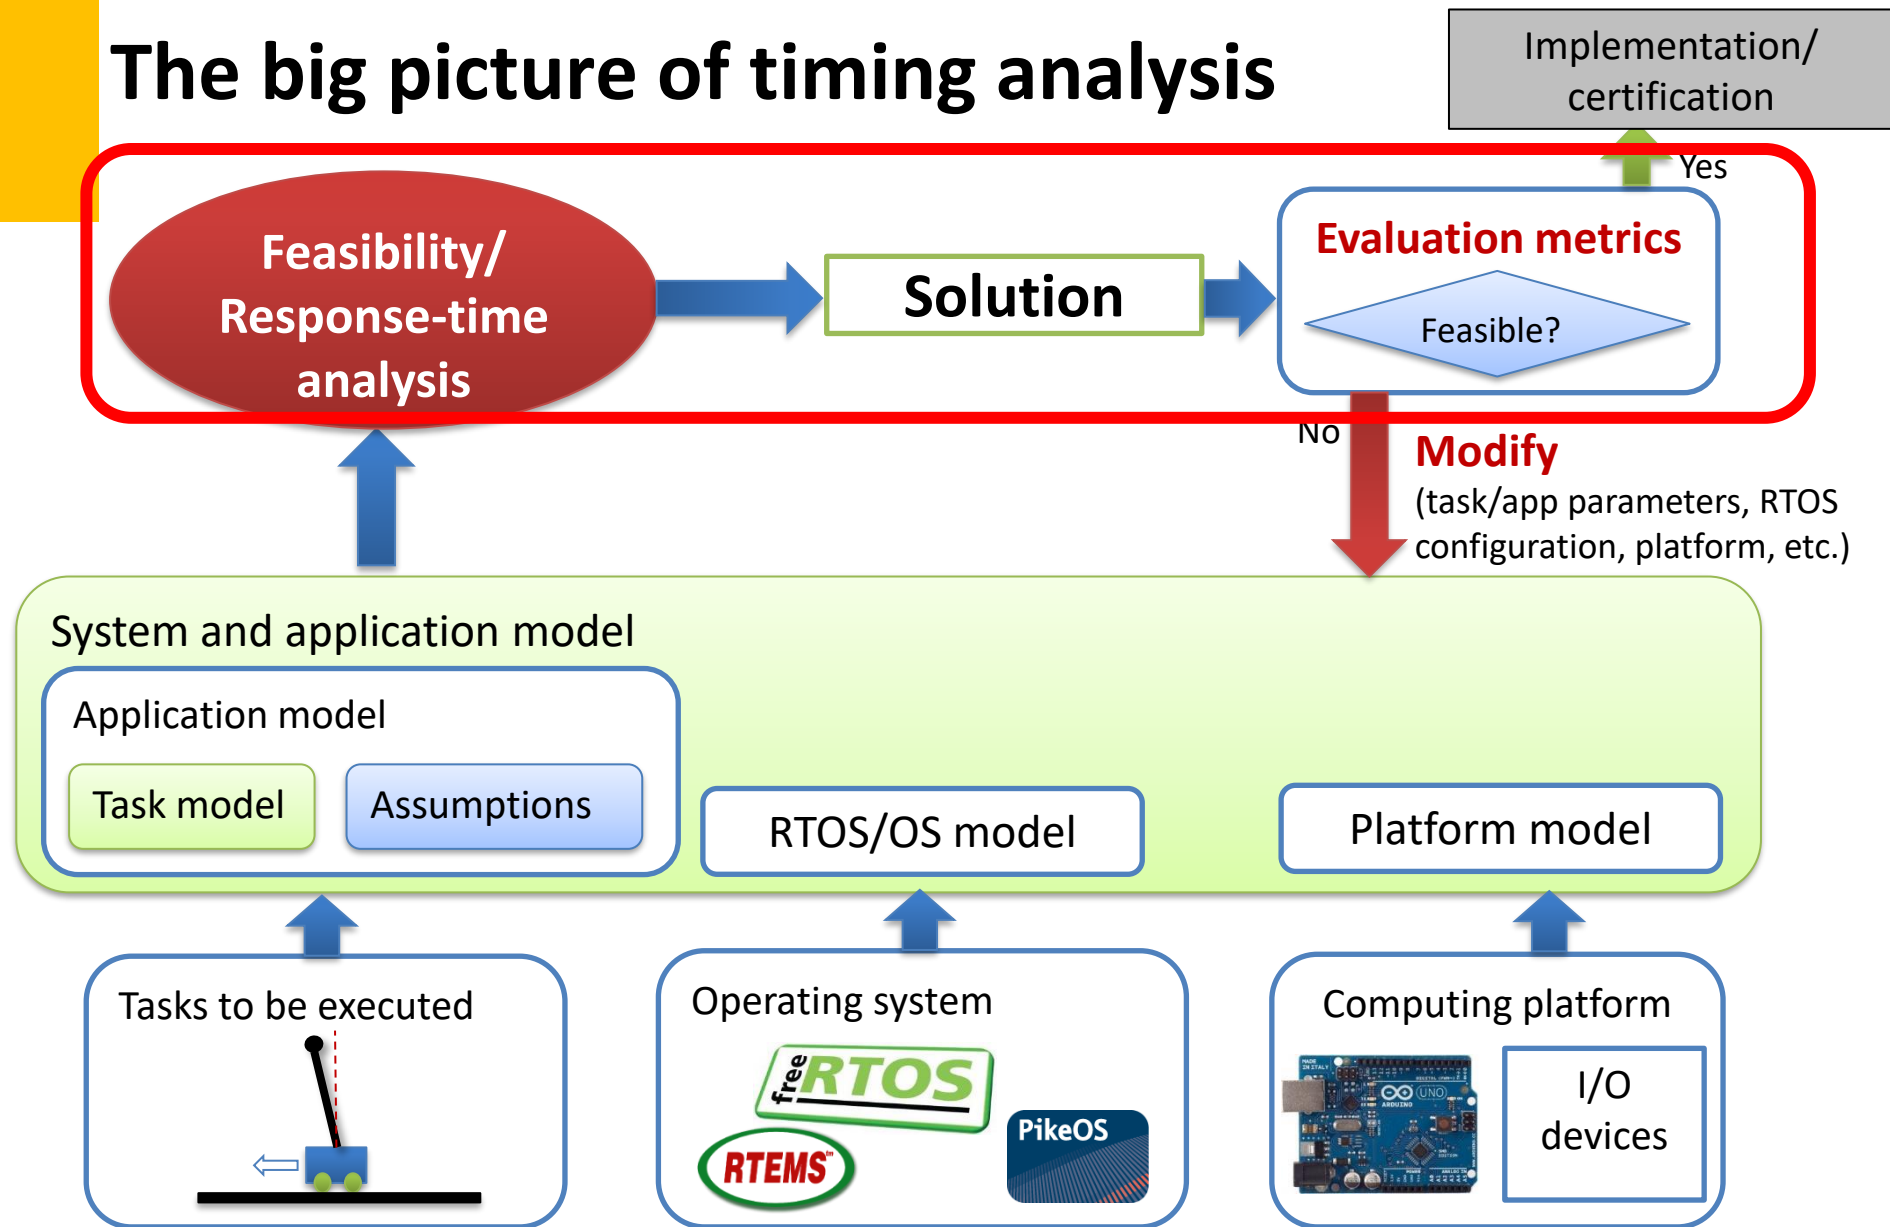
2IMN20 - Real-Time Systems

Schedulability Tests for Periodic (and Sporadic) Tasks

Geoffrey Nelissen


2023-2024

The big picture of timing analysis



Online scheduling of periodic tasks

Buttazzo's book, chapter 4



Disclaimer:

Many slides were provided by Dr. Mitra Nasri

Some slides have been taken from [Giorgio Buttazzo](#)

Agenda

- **Necessary vs sufficient schedulability tests**
- **EDF** schedulability test
- **Priority assignment for task-level fixed-priority scheduling**
 - Rate monotonic (RM)
 - Deadline monotonic (DM)
 - Audsley's Optimal priority assignment algorithm (OPA)
- **RM schedulability tests**
 - Liu and Layland's test [1973]
 - Hyperbolic bound [2000]
 - A utilization-based test for harmonic tasks

Recall: notations

We consider a computing system that has to execute a set τ of n periodic real-time tasks:

$$\tau = \{\tau_1, \tau_2, \tau_3, \dots, \tau_n\}$$


T_i = period

C_i = worst-case execution time (WCET)

D_i = relative deadline

ϕ_i = offset (or phase)

σ_i = release jitter



Definition: utilization

- Each task uses the processor for a fraction of time:

$$U_i = \frac{C_i}{T_i}$$

- Hence the total **processor utilization** is:

$$U = \sum_{i=1}^n \frac{C_i}{T_i}$$

- U is a measure of the **processor load**

Assumptions in this lecture


- We assume
 - A1.** Tasks are fully preemptive
 - A2.** Context switch, preemption, and scheduling overheads are zero
 - A3.** Tasks are independent:
 - no precedence relations
 - no resource constraints
 - No shared resource accesses
 - ...
 - A4.** No self-suspension
 - no blocking on I/O operations
 - ...

Feasibility of a task set

A task set τ is **feasible if and only if** there always exists a schedule in which all tasks meet their timing constraints.


Feasibility of a task set

A task set τ is **feasible if and only if** there always exists a schedule in which each task $\tau_i \in \tau$ can execute for C_i units of time within every interval $[r_{i,k}, d_{i,k})$ for all $k \in \mathbb{N}$.



Feasible schedule


A **schedule** of task set τ is **feasible if and only if** it respects all timing constraints of τ .



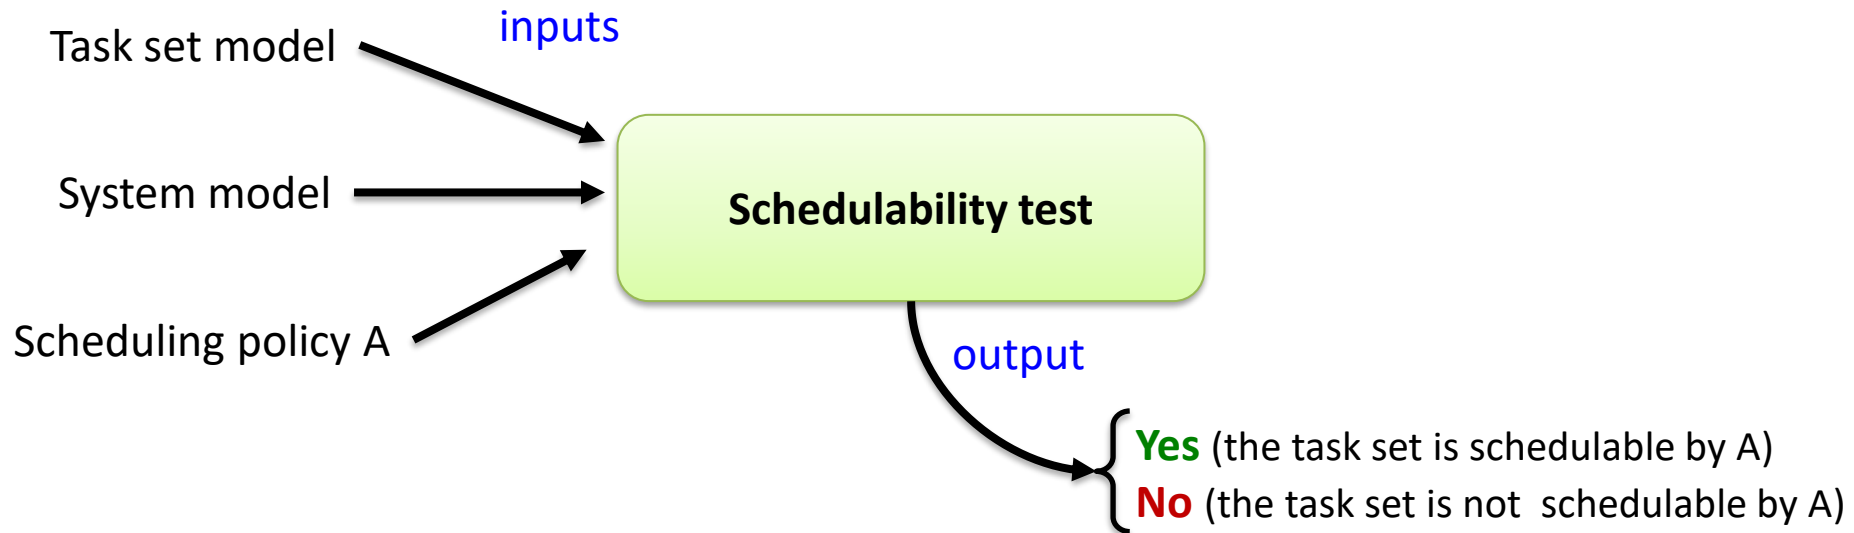
Schedulability of a task set with algorithm A

A task set τ is **schedulable with algorithm A** **if and only if** algorithm A always generates a feasible schedule for τ .

Reminder: Feasibility vs. schedulability



Schedulability test



Necessary vs sufficient **schedulability** tests

Necessary schedulability test for a scheduling algorithm A:
If a task set is **schedulable** by the scheduling algorithm A
then it **certainly passes the test**

Sufficient schedulability test for a scheduling algorithm A:
if a task set that **passes the test**, then it is **certainly**
schedulable by the scheduling algorithm **A**

Exact schedulability test for a scheduling algorithm A:
A task set that **passes the test if and only if** it is **schedulable** by **A**

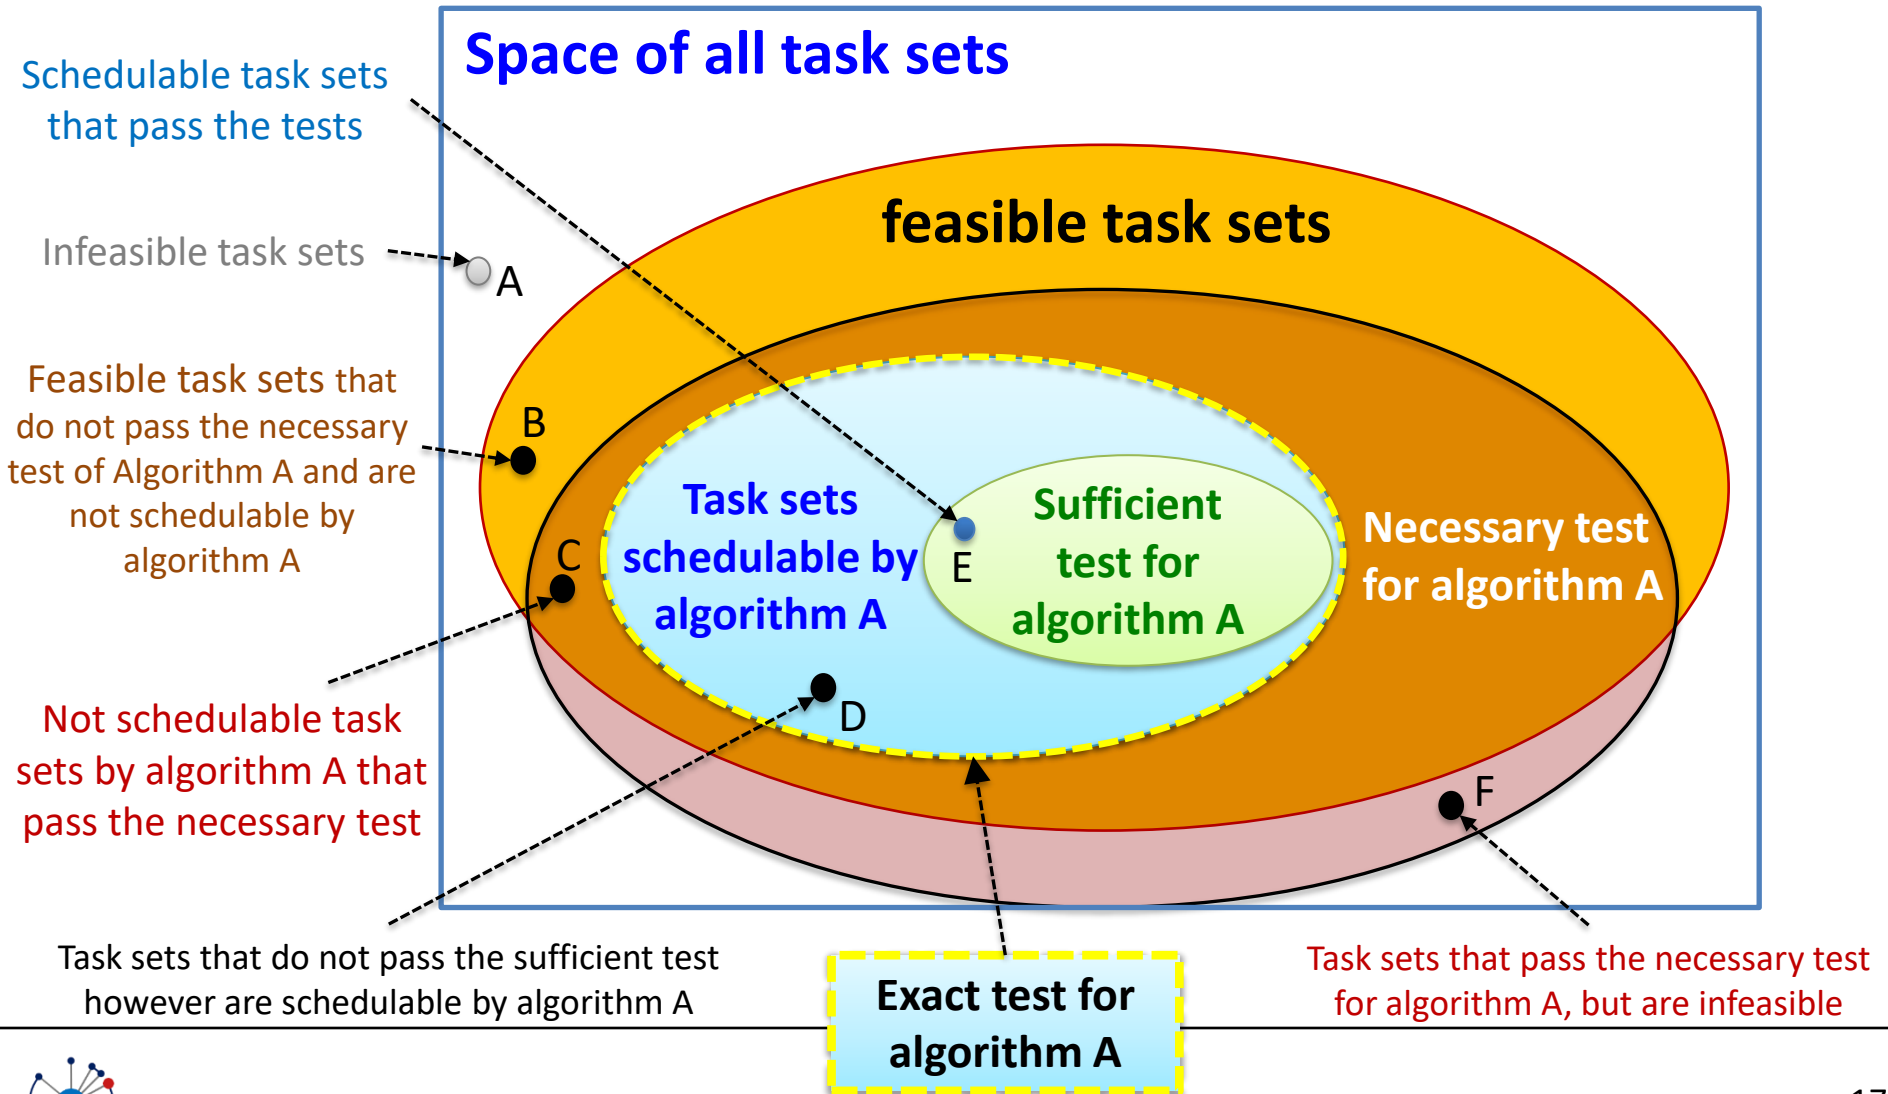
Necessary vs sufficient **schedulability** tests

A **schedulability test** for a scheduling algorithm **A** is a **necessary test** if any task set that **does NOT pass** the test is certainly **not schedulable** by the scheduling algorithm **A**

A **schedulability test** for a scheduling algorithm **A** is a **sufficient test** if any task set that **passes** the test is certainly **schedulable** by the scheduling algorithm **A**

A **schedulability test** for a scheduling algorithm **A** is an **exact test** if any task set that **passes** the test is certainly **schedulable** and any task set that **does not pass** the test is certainly **not schedulable** by the scheduling algorithm **A**

Necessary vs sufficient schedulability tests



Utilization-based schedulability tests

A **utilization-based schedulability test** for an **algorithm A** is a test that checks whether a function f_A on the **utilizations** of the tasks in the task set holds.

A necessary utilization-based schedulability test for uniprocessor scheduling

A necessary test for uniprocessor

$$U = \sum_{i=1}^n \frac{C_i}{T_i}$$

Can we find a feasible schedule for a task set whose utilization is larger than 1? ($U > 1$)

A: yes

C: depends on the task set

B: depends on the scheduling policy

D: no

A necessary test for uniprocessor

$$U = \sum_{i=1}^n \frac{C_i}{T_i}$$

Can we find a feasible schedule for a task set whose utilization is larger than 1? ($U > 1$)

A: yes

C: depends on the task set

B: depends on the scheduling policy

D: no

If $U > 1$, then the amount of work to be done per unit of time **is larger than the unit of time itself!**



A utilization-based schedulability test for EDF

EDF

Schedulability test

- Liu and Layland 1973

A preemptive task set τ with $\forall i, D_i = T_i$, is schedulable by EDF **if and only if**:

$$\sum_{i=1}^n \frac{C_i}{T_i} \leq 1$$

Any feasible periodic or sporadic task set with $\forall i, D_i = T_i$
and $U \leq 1$ can be scheduled by EDF

How is such a task called?

This test is both **necessary and sufficient for EDF** (for implicit deadline periodic or sporadic tasks executed on a single core platform)

EDF Optimality

EDF is optimal among all algorithms:

If there exists a feasible schedule for Γ , then EDF will generate a feasible schedule.



If Γ is not schedulable by EDF, then it cannot be scheduled by any algorithm.

Task-level fixed priority scheduling

Fixed-priority scheduling

Two steps to design a timing-predictable periodic task system under fixed-priority scheduling:

1. **Assign priorities** to each task based on its timing constraints.
 1. Rate monotonic
 2. Deadline monotonic
 3. Optimal priority assignment algorithm (OPA)
2. **Verify the schedulability** of the task set using a **schedulability test**.

In this lecture we limit ourselves to presenting tests for RM. We will cover all priority assignment policies in the next lecture.

Priority assignment for fixed-priority scheduling


Rate monotonic

- Assign priorities monotonically with the activation frequency, a.k.a., *rate* ($\sim 1 / T$) such that a task with a smaller period gets a higher priority
- Example: $T = 10 \Rightarrow \text{rate} = \frac{1}{10}$

τ_i	C_i	T_i	D_i	U_i
τ_1	1	5	5	0.2
τ_2	10	40	15	0.25
τ_3	1	20	20	0.05

Priority ordering?

$P1 > P3 > P2$



Priority assignment for fixed-priority scheduling


Deadline monotonic

- Assign priorities monotonically with the relative deadline of the task, ($\sim 1 / D$) such that a task with a smaller relative deadline gets a higher priority

τ_i	C_i	T_i	D_i	U_i
τ_1	1	5	5	0.2
τ_2	10	40	15	0.25
τ_3	1	20	20	0.05

Priority ordering?

P1 > P2 > P3



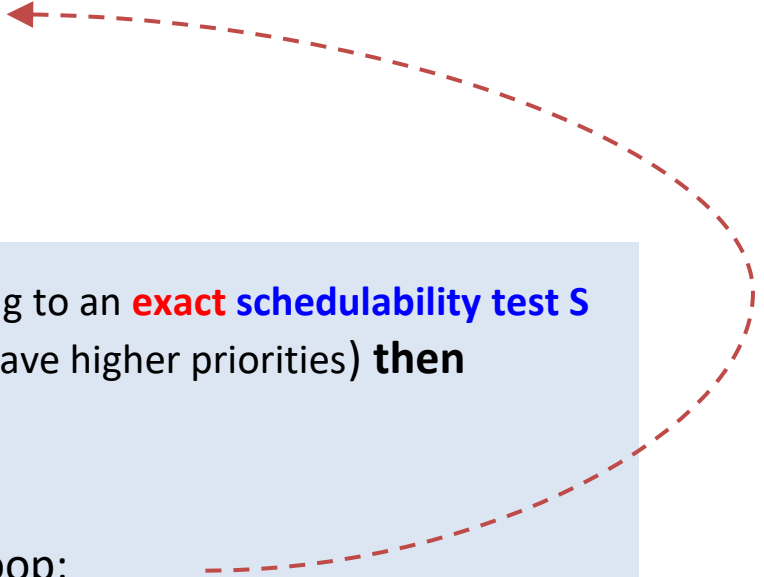
Audsley's **optimal priority assignment** algorithm (OPA)

```
for (each priority level  $k$ , lowest first)
{
    for (each unassigned task  $\tau_i$ )
    {

    }
    return unschedulable;
}
return schedulable (return priorities);
```

Audsley's **optimal priority assignment** algorithm (OPA)

```
for (each priority level  $k$ , lowest first)
{
  for (each unassigned task  $\tau_i$ )
  {
    if ( $\tau_i$  is schedulable at priority  $k$  according to an exact schedulability test S
        with all unassigned tasks assumed to have higher priorities) then
    {
      assign  $\tau_i$  to priority  $k$ ;
      break and continue the outer loop;
    }
  }
  return unschedulable;
}
return schedulable (return priorities);
```



Is FP an optimal scheduling policy with RM, DM or Audsley's priority assignment?

What do you think?

The limitation of fixed-priority scheduling

Fixed-priority scheduling with **RM** priorities is **not** an **optimal** scheduling policy (in the sense of feasibility)


Neither is it with DM or Audsley's priority assignment

RM-priority assignment is not optimal

Build a feasible task set with $U \leq 1$ that is not schedulable by RM (rate monotonic)

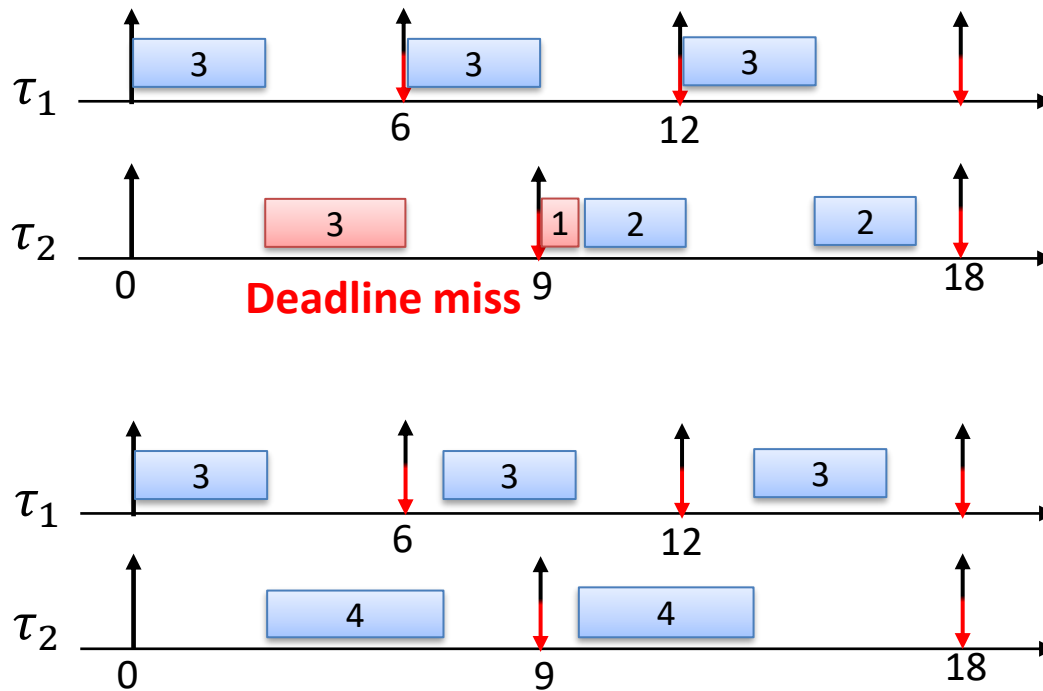
Hint:

What execution times will lead to $U < 1$ and will result in a deadline miss for one of the tasks?



RM-priority assignment is not optimal

Build a feasible task set with $U \leq 1$ that is not schedulable by RM (rate monotonic)



τ_i	C_i	T_i	D_i	U_i
τ_1	3	6	6	0.5
τ_2	4	9	9	0.44

$$U = \sum_{i=1}^n U_i = 0.5 + 0.44 = 0.94$$

A feasible schedule

Fixed-priority scheduling with RM priorities is **not** an **optimal** scheduling policy (in the sense of feasibility)

However, if $\forall i, D_i = T_i$ then
the **rate-monotonic priority assignment** is
an **optimal priority assignment** among all *other fixed* priority assignments

Rate Monotonic is optimal

RM is **optimal** among all fixed priority algorithms (if $D_i = T_i$):

If there exists a fixed priority assignment which leads to a feasible schedule for Γ , then the RM assignment is feasible for Γ .



If Γ is not schedulable by RM, then it cannot be scheduled by any fixed priority assignment.


How is such a task called?

Deadline Monotonic is optimal

If $D_i \leq T_i$ then the **optimal** priority assignment is given by **Deadline Monotonic (DM)**:

DM

$P2 > P1$

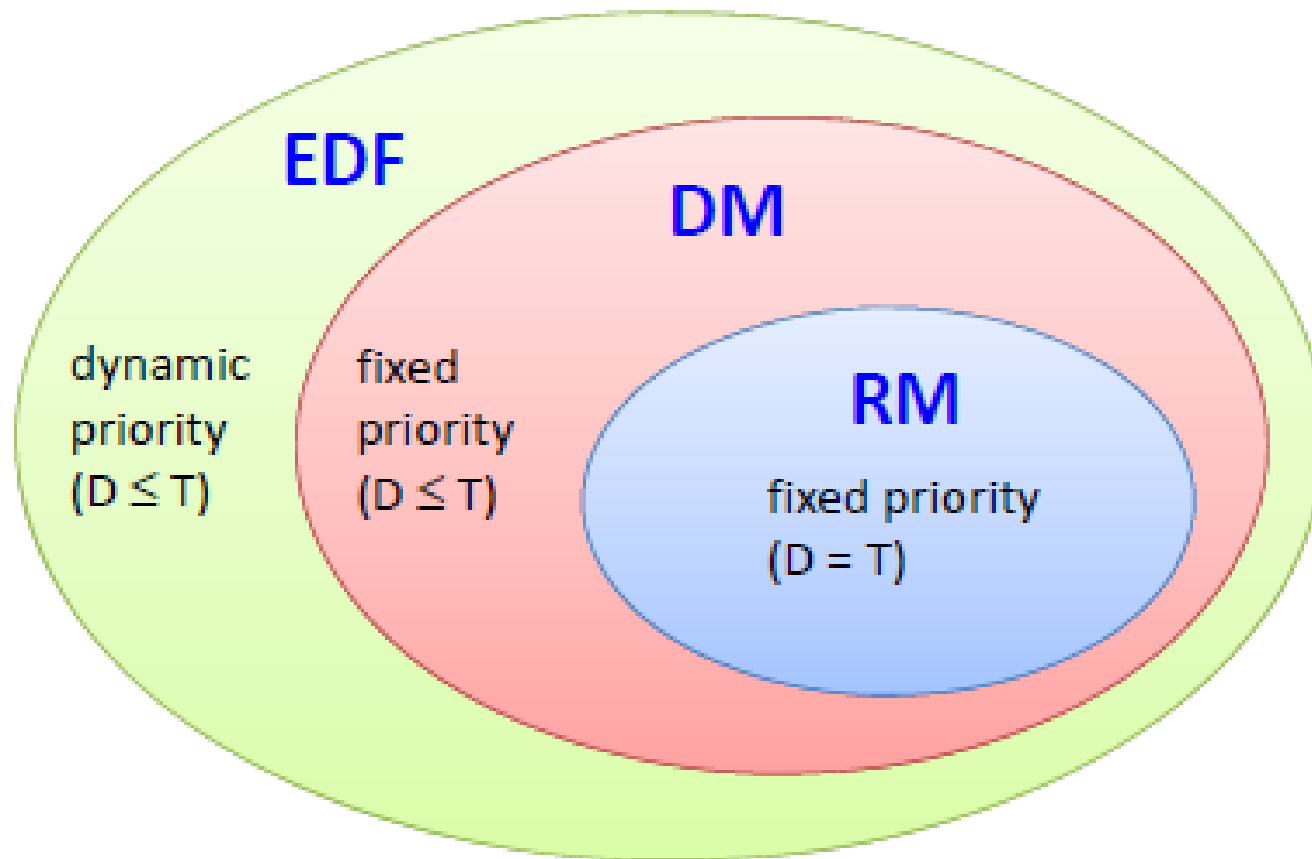


RM

$P1 > P2$



Optimality



Utilization-based schedulability tests **for RM**

Building a first utilization-based test for RM

Find a utilization **threshold** (i.e, a bound) such that **ANY** **task set** with **utilization lower than that bound** is **CERTAINLY schedulable by RM**



The simplest utilization-based test

For a given task set, check whether or not

$$U \leq U_{th}$$

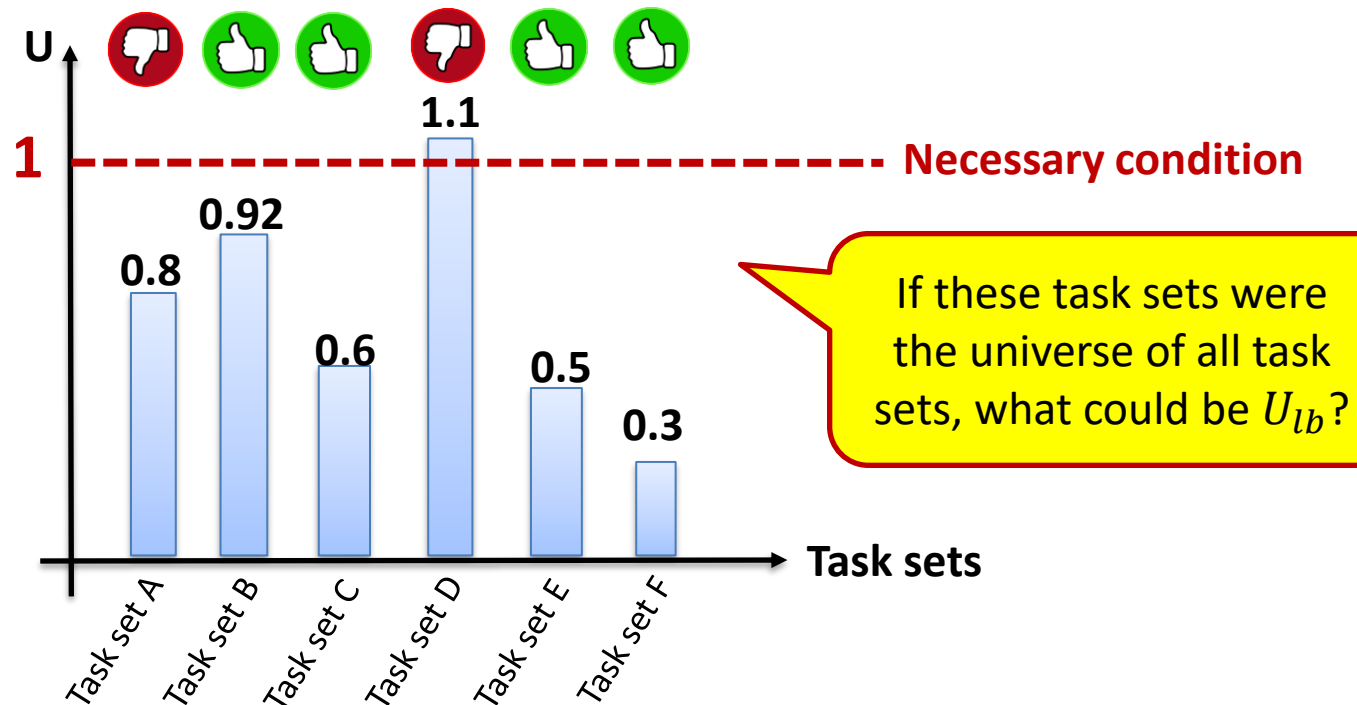
where U_{th} is the largest utilization such that any task set with $U \leq U_{th}$ is **always schedulable by RM**.

The simplest utilization-based test

For a given task set, check whether or not

$$U \leq U_{th}$$

where U_{th} is the largest utilization such that any task set with $U \leq U_{th}$ is **always schedulable by RM**.




The simplest utilization-based test


For a given task set, check whether or not

$$U \leq U_{th}$$

where U_{th} is the largest utilization such that any task set with $U \leq U_{th}$ is **always schedulable by RM**.




Schedulability of a FP policy is not monotonic with the utilization!



τ_i	C_i	T_i	D_i	U_i
τ_1	3	6	6	0.5
τ_2	4	9	9	0.44

$$U = 0.94$$

$U = 0.94$ (smaller than 1) but not schedulable!



τ_i	C_i	T_i	D_i	U_i
τ_1	3	6	6	0.5
τ_2	6	12	12	0.5

$$U = 1$$

$U = 1$, but it is schedulable by RM!

In this task set, periods are harmonic. Each period divides all smaller ones.

Liu and Layland test for RM

Liu and Layland [1973] derived a value for U_{lb} for the **rate monotonic** scheduling under certain **assumptions**:

- A1.** Every job of τ_i executes for its WCET C_i
- A2.** For each task, $T_i = D_i$
- A3.** Tasks are fully preemptive
- A4.** Context switch, preemption, and scheduling overheads are zero
- A5.** Tasks are sequential and independent:
 - no precedence relations
 - no resource constraints
 - no blocking on I/O operations
 - no self suspension
 - No shared resource accesses

Hint for the exam: remember the assumptions

Liu and Layland's test for RM

n is the number of tasks in the task set

$$U \leq n \cdot (2^{1/n} - 1)$$

$$\begin{array}{l} n \rightarrow \infty \\ U_{lb} \rightarrow ? \end{array}$$




$$\begin{array}{l} n \rightarrow \infty \\ U_{lb} \rightarrow \ln 2 \sim 0.691 \end{array}$$

$$\begin{array}{l} n \rightarrow 2 \\ U_{lb} \rightarrow ? \end{array}$$



$$\begin{array}{l} n \rightarrow 2 \\ U_{lb} \rightarrow 0.83 \end{array}$$


Liu and Layland (L&L) test for RM



Liu and Layland's test accepts more task sets when n is small


How happy are we with the L&L RM test?

$U \leq 1$ is necessary and sufficient
for EDF schedulability



$$U = \sum_{i=1}^2 U_i = U_1 + U_2$$

Hyperbolic bound



The Hyperbolic Bound

- In 2000, **Bini et al.** proved that a set of n periodic tasks is schedulable with RM if:

$$\prod_{i=1}^n (U_i + 1) \leq 2$$

Example

Is the task set feasible?

Yes $U \leq 1$

Does the task set pass the Liu & Layland's test?

No because $U = 0.85 > 2(2^{1/2} - 1) \sim 0.83$

Does the task set pass the hyperbolic-bound test?

Yes because $(0.8 + 1) \times (0.05 + 1) = 1.89 < 2$

Is the task set schedulable by RM?

Yes! Because it passed the hyperbolic bound test!

τ_i	C_i	T_i	D_i	U_i
τ_1	8	10	10	0.8
τ_2	0.9	18	18	0.05

$$U = 0.85$$

$$\sum_{i=1}^n U_i \leq 1$$

necessary

$$\sum_{i=1}^n U_i \leq n(2^{1/n} - 1)$$

Liu and Layland test

$$\prod_{i=1}^n (U_i + 1) \leq 2$$

Hyperbolic bound

Example

Is the task set feasible?

Yes $U \leq 1$

Does the task set pass the Liu & Layland's test?

No because $U > 2(2^{1/2} - 1) \sim 0.83$

Does the task set pass the hyperbolic-bound test?

No because $(0.4 + 1) \times (0.6 + 1) = 2.24 > 2$

Is the task set schedulable by RM?

Inconclusive. Those utilization bounds are **only sufficient tests**, they are **not exact**.

(Note: the task set is in fact schedulable with RM)

$$\sum_{i=1}^n U_i \leq 1$$

necessary

$$\sum_{i=1}^n U_i \leq n(2^{1/n} - 1)$$

Liu and Layland test

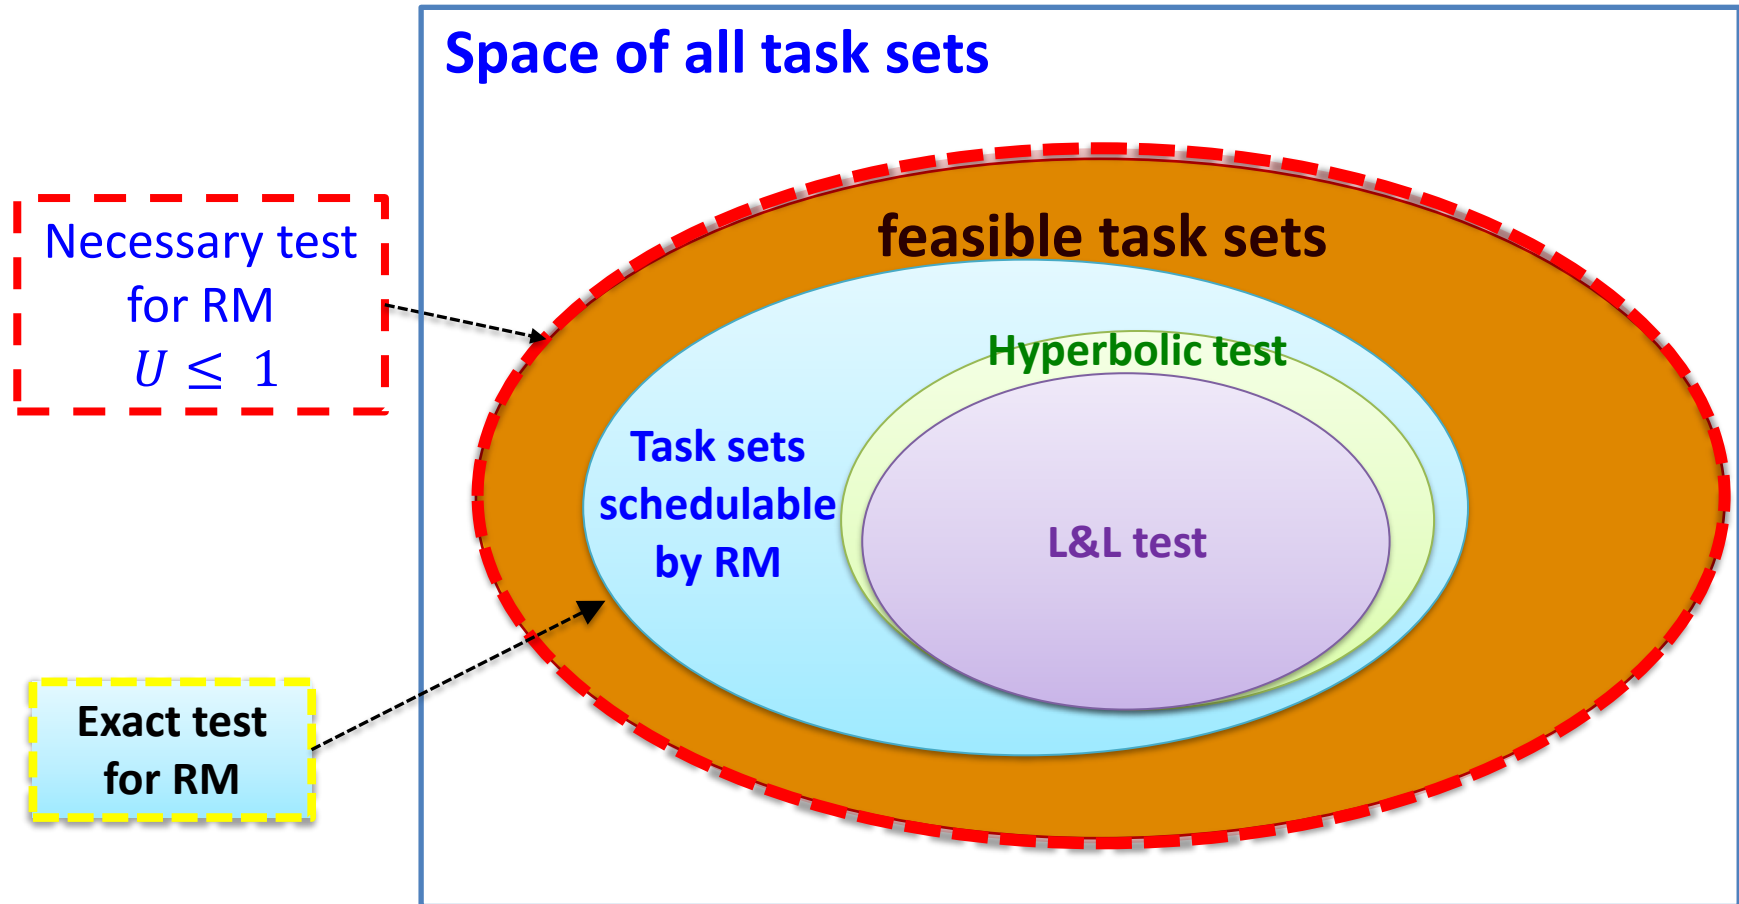
$$\prod_{i=1}^n (U_i + 1) \leq 2$$

Hyperbolic bound

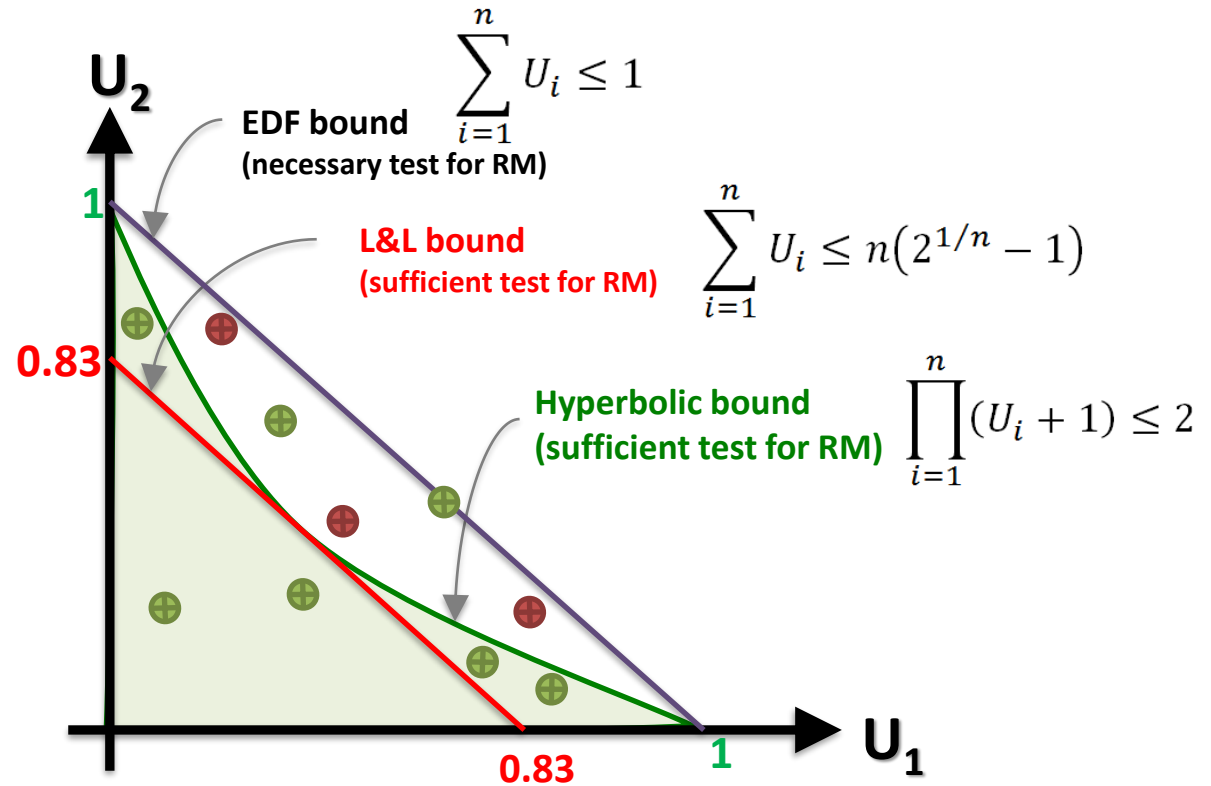
τ_i	C_i	T_i	D_i	U_i
τ_1	3	5	5	0.6
τ_2	4	10	10	0.4

$$U = 1$$

The universe of utilization-based tests for RM



Utilization-based tests and RM scheduling



Bad news:

We cannot have a better utilization-based test than the hyperbolic bound!



Hyperbolic bound is tight

It is **impossible** to build a new **utilization-based test A** such that it accepts more task sets than the hyperbolic bound test!

In other words, as long as the only information used in a test is the **tasks' utilization**, that test will **never be better** than the hyperbolic bound test!

An improved test for **harmonic tasks** scheduled under RM

A utilization-based test for harmonic tasks

Han et al. [1997] have proven that **rate monotonic** is **optimal** if $\forall \tau_i, D_i = T_i, \sigma_i = 0$ and periods are **harmonic** (task periods are multiples of each others).

Hence, if periods are **harmonic** and $\forall \tau_i, D_i = T_i$, the following schedulability test is both **necessary and sufficient** for RM

$$U \leq 1$$

Try to explain why

C.-C. Han and H.-Y. Tyan. A Better Polynomial-time Schedulability Test for Real-time Fixed-priority Scheduling Algorithms. In IEEE Real-Time Systems Symposium (RTSS), pages 36–45, 1997.



QUIZ TIME

Quiz

- Is this task set feasible?
- Yes
- No
- Maybe (depends on the scheduling algorithm)

τ_i	C_i	T_i	D_i	U_i
1	1	5	5	0.2
2	5	10	10	0.5
3	1	10	10	0.1
4	1	10	10	0.1
5	1	15	15	0.06

$$U = 0.96$$

Quiz

- Is this task set feasible?
- Yes
- No
- Maybe (depends on the scheduling algorithm)

τ_i	C_i	T_i	D_i	U_i
1	1	5	5	0.2
2	5	10	10	0.5
3	1	10	10	0.1
4	1	10	10	0.1
5	1	15	15	0.06

$$U = 0.96$$

Quiz

- Based on its utilization, is this task set schedulable with RM?
- Yes
- No
- Maybe

τ_i	C_i	T_i	D_i	U_i
1	1	5	5	0.2
2	5	10	10	0.5
3	1	10	10	0.1
4	1	10	10	0.1
5	1	15	15	0.06

$$U = 0.96$$

Quiz

- Based on its utilization, is this task set schedulable with RM?
- Yes
- No
- Maybe

τ_i	C_i	T_i	D_i	U_i
1	1	5	5	0.2
2	5	10	10	0.5
3	1	10	10	0.1
4	1	10	10	0.1
5	1	15	15	0.06

$$U = 0.96$$

What defines an optimal scheduling policy (in the sense of feasibility)?

- For infeasible task sets, it minimizes the number of deadline misses
- It always generates a feasible schedule
- It generates a feasible schedule for a feasible task set

What defines an optimal scheduling policy (in the sense of feasibility)?

- For infeasible task sets, it minimizes the number of deadline misses
- It always generates a feasible schedule
- It generates a feasible schedule for a feasible task set

Which policy is optimal for preemptive independent tasks on single core?

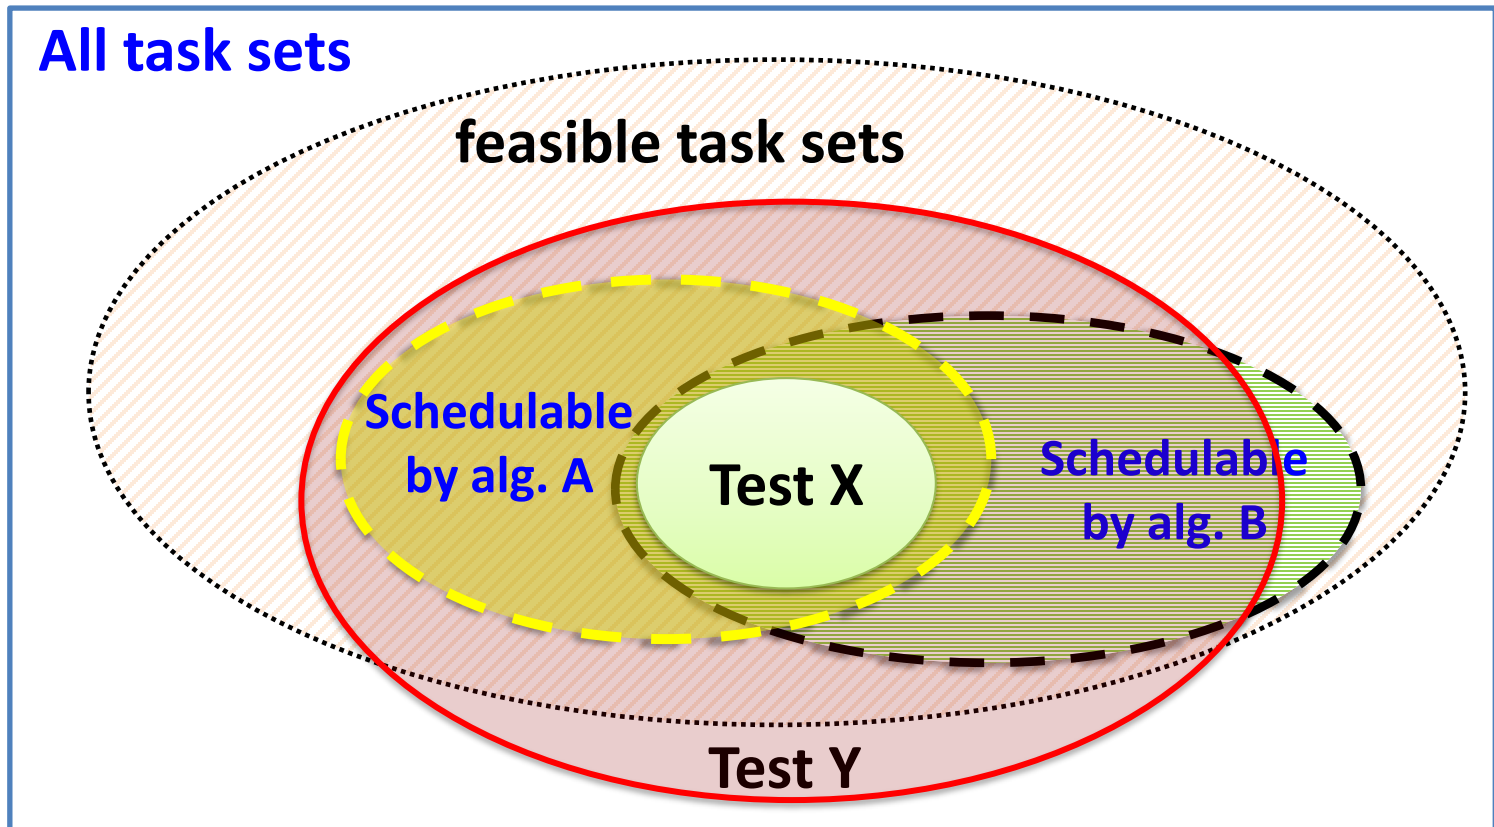
- RM
- FIFO
- EDF
- DM

Which policy is optimal for preemptive independent tasks on single core?

- RM
- FIFO
-  • EDF
- DM

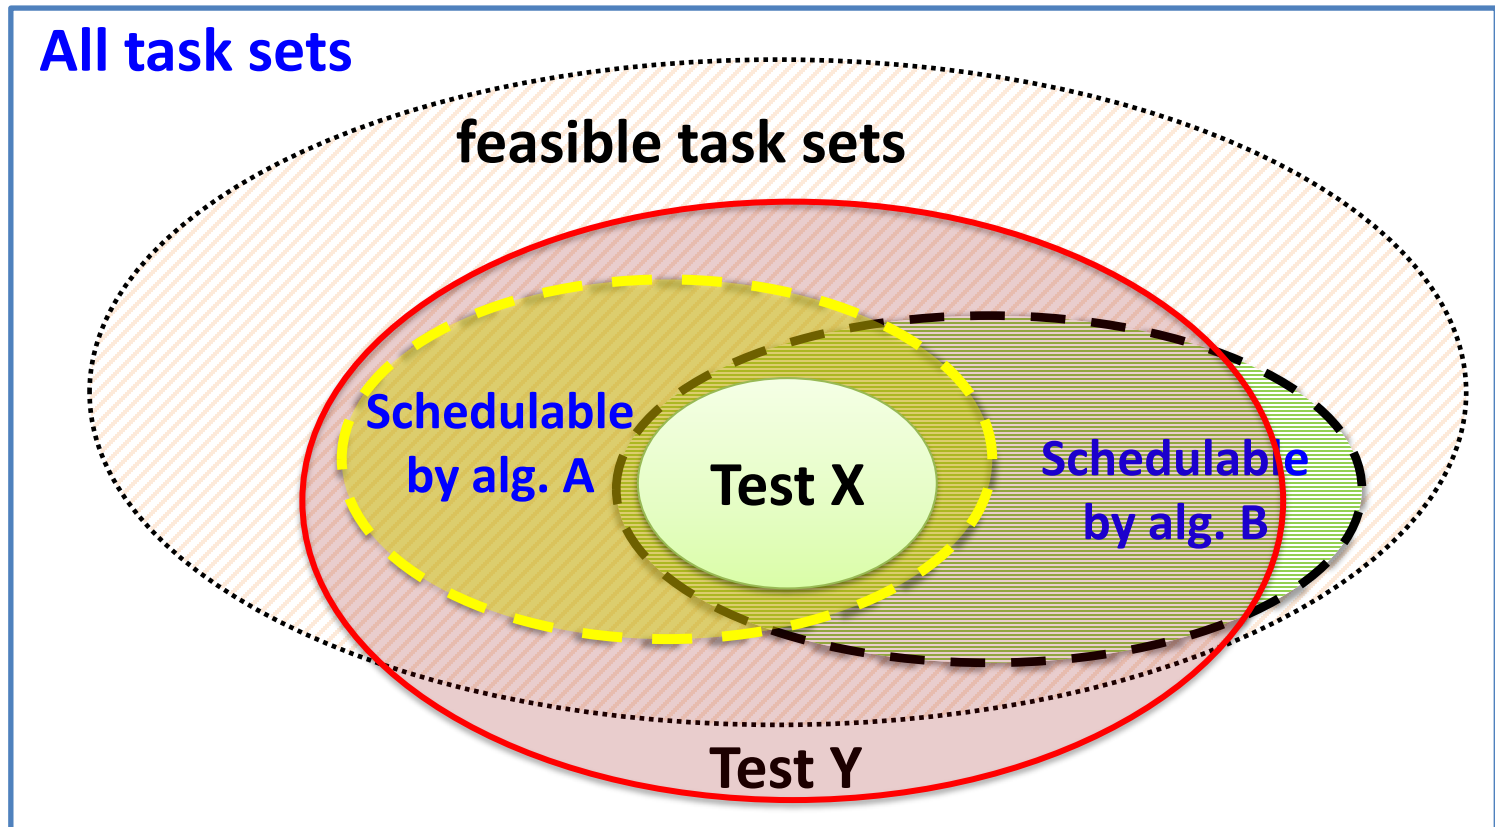
Pick a correct answer

All task sets



- Test X is a necessary test for alg. A
- An exact test for B is a sufficient test for A
- Test X is a sufficient test for A and B
- Test Y is not a necessary test for any algorithm

Pick a correct answer



- Test X is a necessary test for alg. A
- An exact test for B is a sufficient test for A
- **Test X is a sufficient test for A and B**
- Test Y is not a necessary test for any algorithm

Summary

- RM, DM, OPA **are all optimal priority assignments** for task-level fixed priority scheduling but under **different sets of assumptions**:
 - **RM**: independent preemptive sporadic or periodic tasks (with $\sigma_i = 0$), single core, $D_i = T_i$ for all tasks
 - **DM**: independent preemptive sporadic or periodic tasks (with $\sigma_i = 0$), single core, $D_i \leq T_i$ for all tasks
 - **OPA**: optimal for any set of preemptive tasks if there exists an **exact schedulability test**
- **EDF** is optimal for independent preemptive tasks on single core
- **Utilization-based** tests for RM
 - **LL-bound**
 - **Hyperbolic bound** (best possible utilization-based test for RM)

**Is rate-monotonic an optimal scheduling policy
(in the sense of feasibility)?**