

2IMN20 - Real-Time Systems


Response-time analysis for fixed-priority preemptive scheduling

Geoffrey Nelissen

2023-2024

Reference book

Chapter 4



Disclaimer:

Many slides were provided by Dr. Mitra Nasri

Some slides have been taken from [Giorgio Buttazzo](#)'s course



Agenda

- RM utilization-based tests (reminder)
- ***Response-time analysis*** for periodic or sporadic tasks under preemptive fixed priority scheduling

Assumptions

- For this lecture, we assume
 - A1.** All jobs of τ_i execute for no more than C_i
 - A2.** Tasks are periodic or sporadic
 - A3.** Tasks are fully preemptive
 - A4.** Context switch, preemption, and scheduling overheads are zero
 - A5.** Tasks are independent:
 - no precedence relations
 - no resource constraints
 - no blocking on I/O operations
 - A6.** No self-suspensions
 - A7.** Single core
 - A8.** **Unknown phasing/offset** for the tasks

Assume that tasks are indexed according to their priority ordering, namely,

$$P_1 > P_2 > P_3 > \dots > P_n$$

Agenda

- RM utilization-based tests (reminder)
- ***Response-time analysis*** for periodic or sporadic tasks under preemptive fixed priority scheduling

Building a utilization-based test for RM

Task utilization

$$U_i = \frac{C_i}{T_i}$$

Task set utilization

$$U = \sum_{i=1}^n \frac{C_i}{T_i}$$

Finding a utilization **threshold** (lower bound) such that **ANY task set** with **utilization lower than that bound** is **CERTAINLY schedulable by RM**



RM utilization-based tests

Assumptions: no release jitter (i.e., $\sigma_i = 0$) and tasks have implicit deadlines (i.e., $D_i = T_i$)

Liu and Layland bound:
$$\sum_{i=1}^n U_i \leq n(2^{1/n} - 1)$$

Hyperbolic bound:
$$\prod_{i=1}^n (U_i + 1) \leq 2$$



QUIZ TIME



QUIZ TIME

Quiz

- Is this task set schedulable by RM according to the Liu and Layland test?

- Yes
- No

τ_i	C_i	T_i	D_i	U_i
τ_1	1.5	5	5	0.3
τ_2	4	10	10	0.4

$$U = 0.7$$

$$\sum_{i=1}^n U_i \leq n(2^{1/n} - 1)$$

Quiz

- If the hyperbolic bound test accepts a task set, Liu and Layland's test also accepts that task set.
- true
- false

Quiz

- Is this task set schedulable by RM according to the hyperbolic bound test?

- Yes
- No

τ_i	C_i	T_i	D_i	U_i
τ_1	1.5	5	5	0.3
τ_2	4	10	10	0.4

$$U = 0.7$$

Quiz

- The hyperbolic bound test is a necessary schedulability test.
- true
- false

Quiz

- If the Liu and Layland test accepts a task set, then the hyperbolic bound test also accepts that task set.
- true
- false

Quiz

- Both the hyperbolic bound and Liu and Layland tests are sufficient schedulability tests.
- true
- false



QUIZ TIME

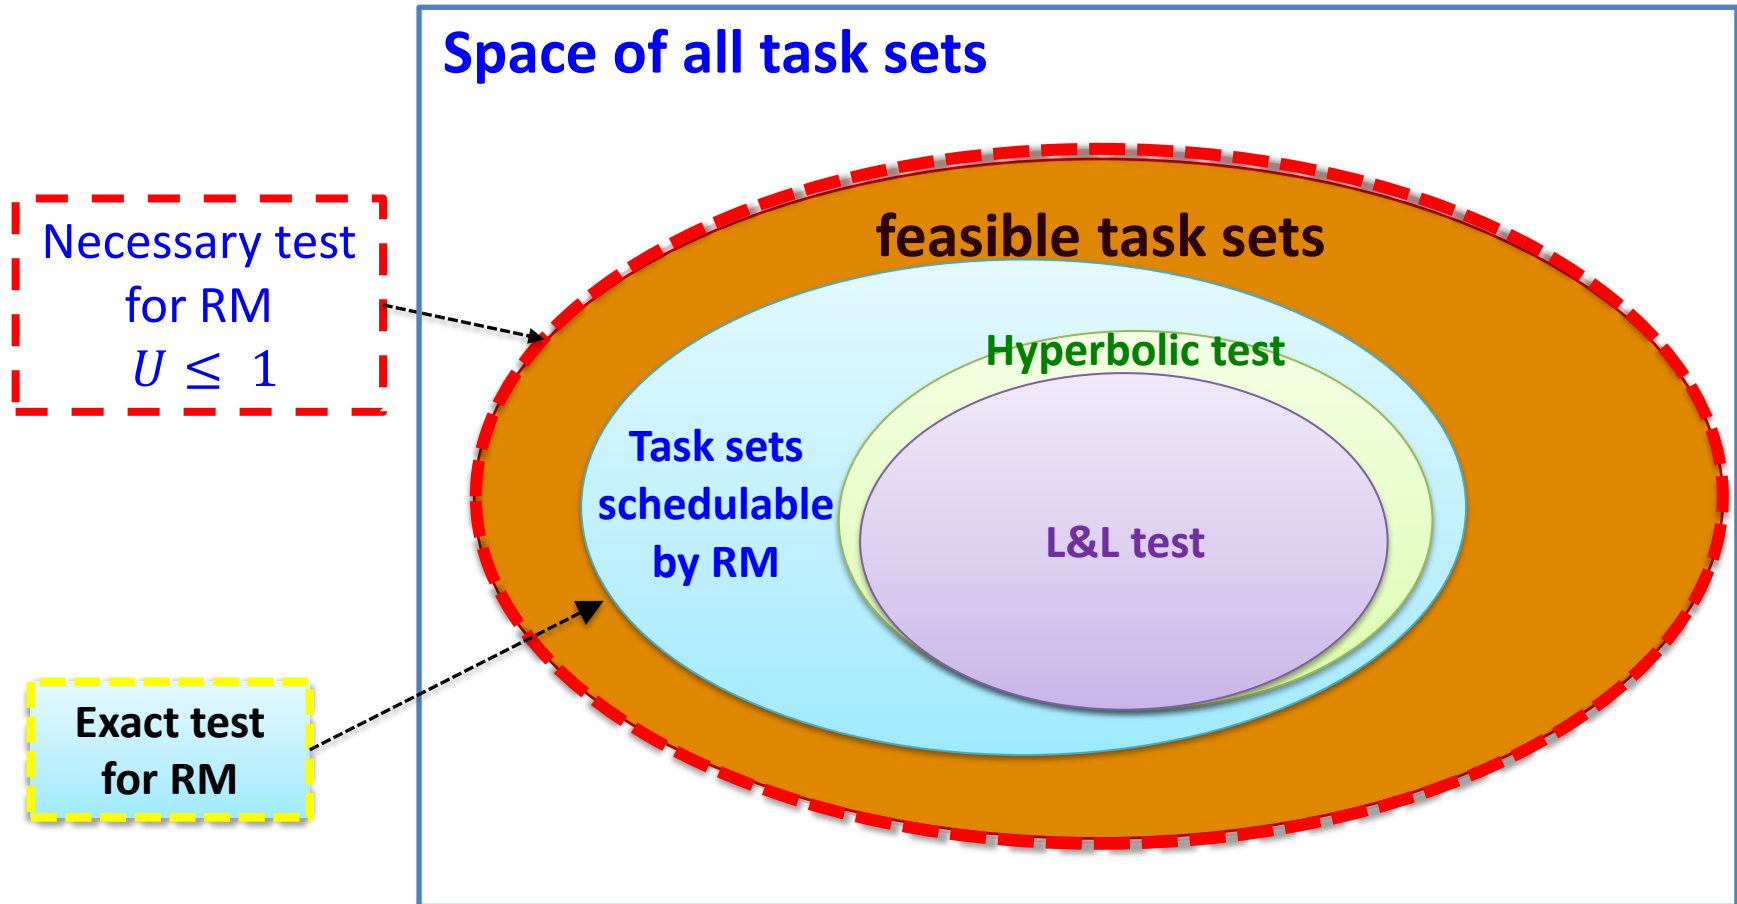
Bonus question

Quiz


- What is the necessary schedulability/feasibility test for single core preemptive scheduling?

$$U \leq 1$$

The universe of utilization-based tests for RM



Previously on Real-Time Systems



Now we introduce
an exact
schedulability test
for RM
(and FP in general)


Agenda

- RM utilization-based tests (reminder)
- *Response-time analysis* for periodic or sporadic tasks under FP
 - **Response-time analysis**
 - For $\forall \tau_i, \sigma_i = 0$ and $R_i \leq T_i$
 - For $\forall \tau_i, \sigma_i \geq 0$ and $R_i \leq T_i$
 - For $\forall \tau_i, \sigma_i \geq 0$ and $R_i \leq T_i$ or $R_i > T_i$
 - **Park's test**

Response-time analysis (RTA)

Unlike utilization-based tests, response-time analysis takes **task's inter-arrival time** and **worst-case execution time (WCET)** into account!

Visualization:



Worst-case response time (WCRT):

$$R_i = \sup\{R_{i,j} \mid \forall j, 1 \leq j \leq \infty\}$$

The schedulability test:

$$\forall \tau_i \in \tau, \quad R_i \leq D_i$$

It is not practical as it is!
We need a smarter solution


A few questions before we start

Under preemptive FP, which tasks **interfere** with the execution of a task τ_i ?

Higher priority tasks

What is the **maximum number of jobs** that a **periodic** task τ_i may release in an **interval of duration L** ?

At most $\left\lceil \frac{L}{T_i} \right\rceil$ jobs if $\sigma_i = 0$



A few questions before we start

Under preemptive FP, which tasks **interfere** with the execution of a task τ_i ?

Higher priority tasks

What is the **maximum number of jobs** that a **periodic** task τ_i may release in an **interval of duration L** ?

At most $\left\lceil \frac{L}{T_i} \right\rceil$ jobs if $\sigma_i = 0$

What is the **maximum number of jobs** that a **sporadic** task τ_i may release in an **interval of duration L** ?

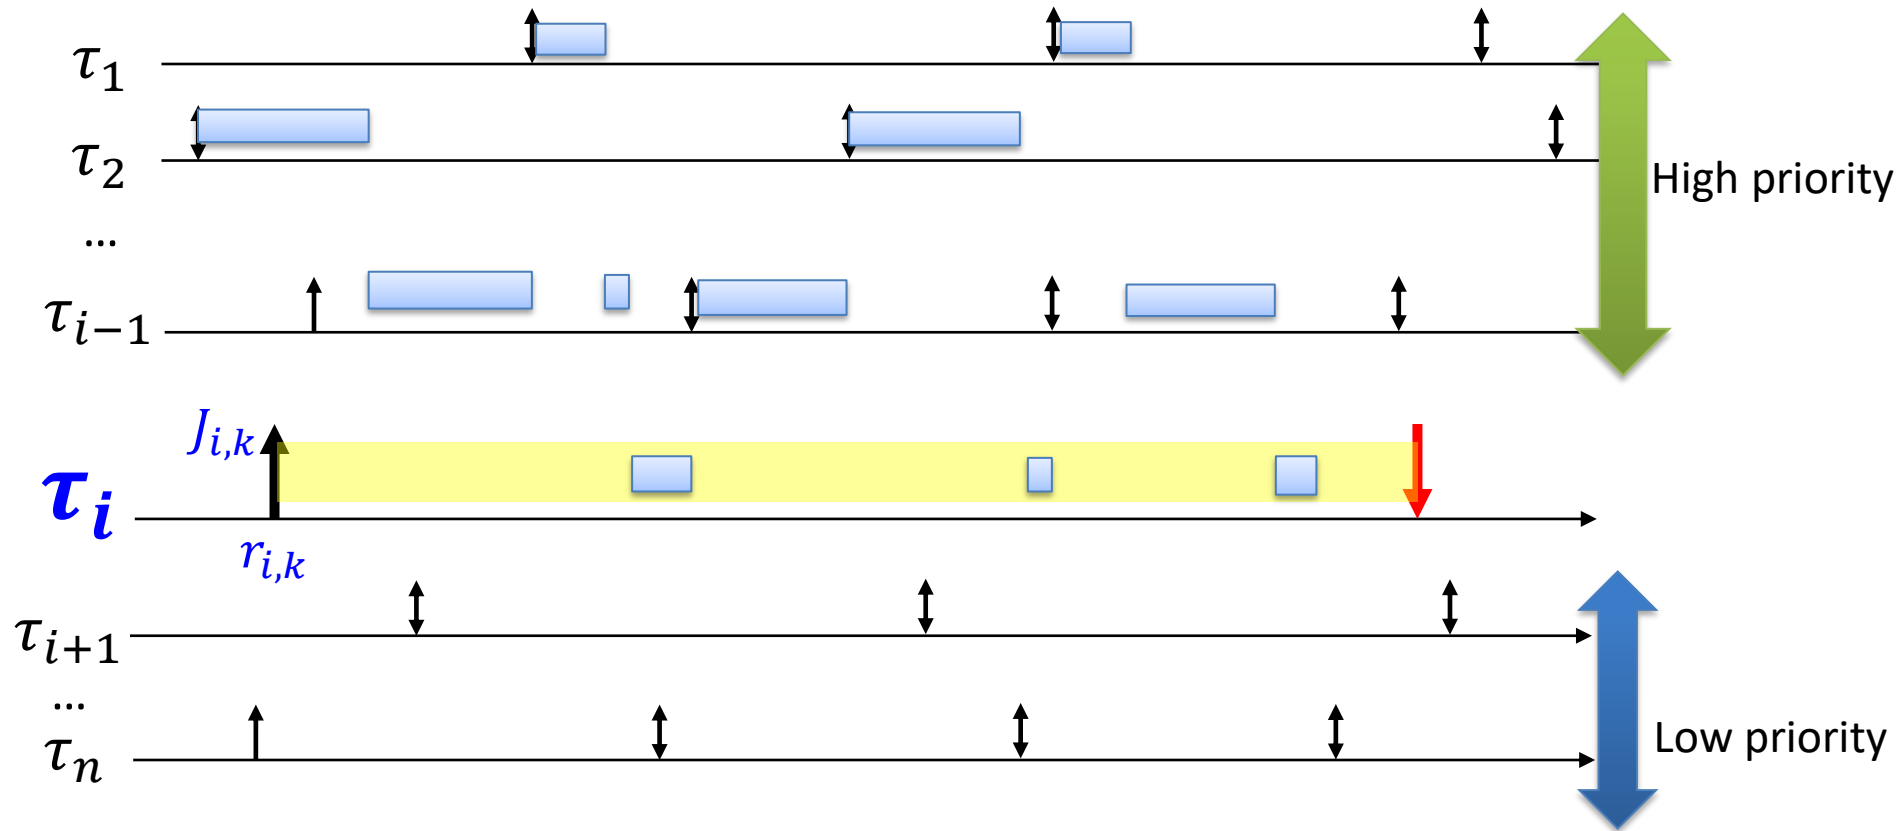
At most $\left\lceil \frac{L}{T_i} \right\rceil$ jobs if $\sigma_i = 0$

What is the **maximum workload** that a periodic or sporadic task τ_i may release in an interval of duration L ?

At most $\left\lceil \frac{L}{T_i} \right\rceil \times C_i$ if $\sigma_i = 0$

On a search for the worst-case release scenario

What is the worst-case release scenario for a job of task τ_i under FP scheduling?




Assume that tasks are sorted and indexed by priority: $P_1 > P_2 > \dots > P_n$

On a search for the worst-case release scenario

The release time of τ_1 is moved to an earlier time to align with τ_i 's release

Response time of τ_i increased!

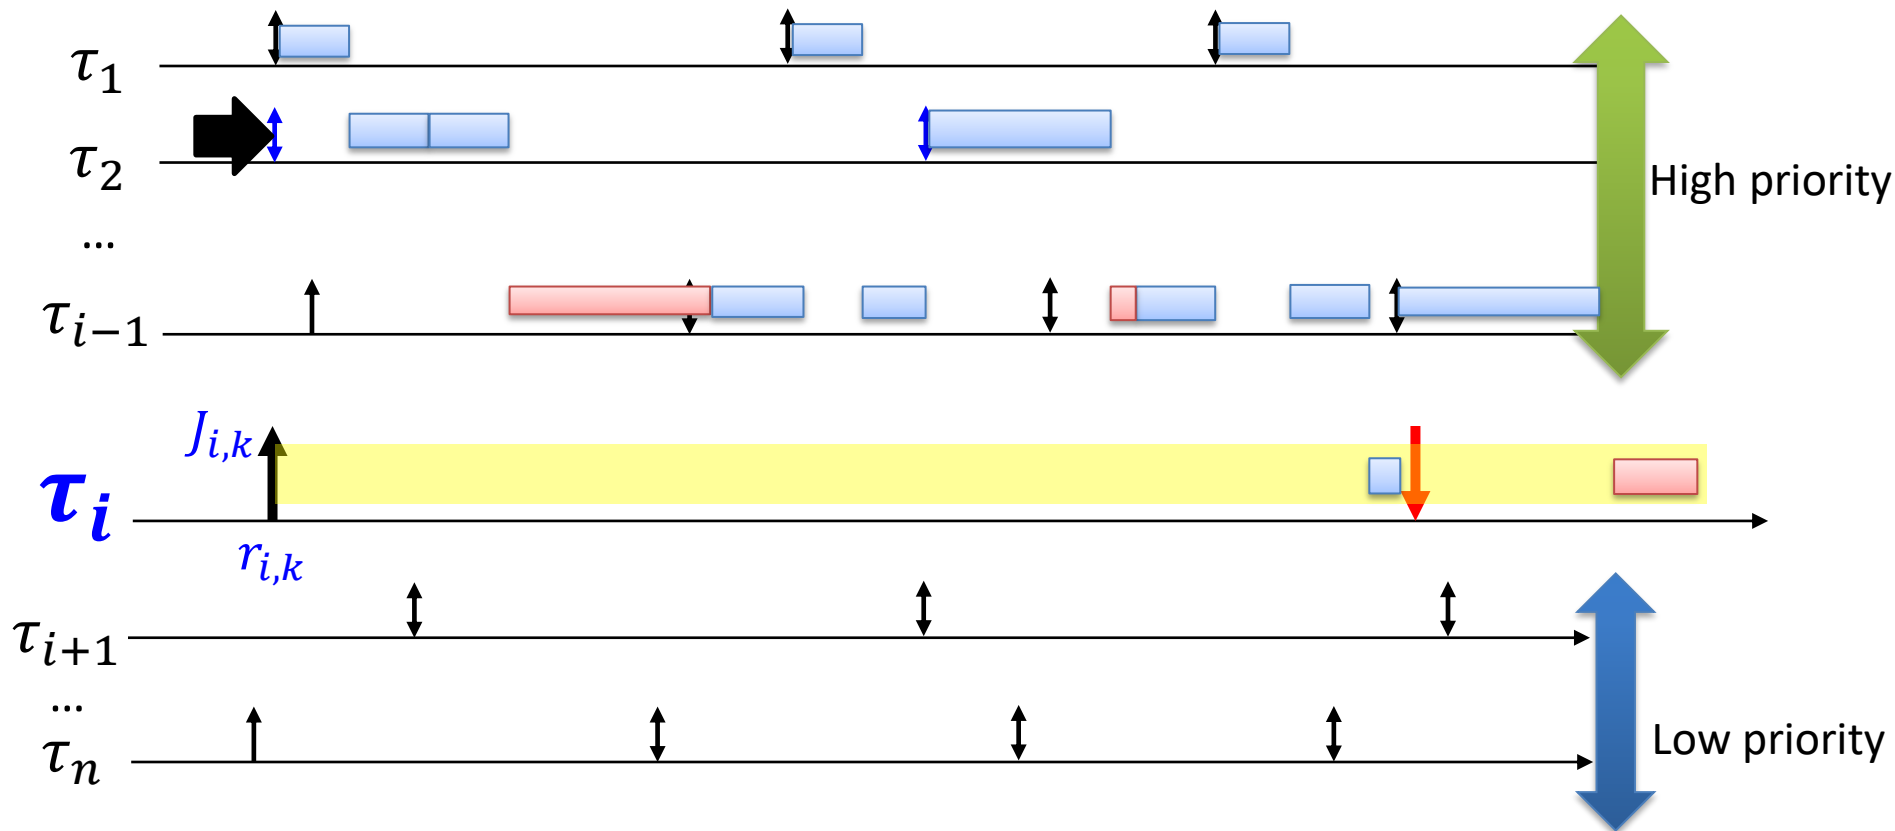


Assume that tasks are sorted and indexed by priority: $P_1 < P_2 < \dots < P_n$

On a search for the worst-case release scenario


The release time of τ_2 is moved to a later time to align with τ_i 's release

Response time of τ_i increased!




Assume that tasks are sorted and indexed by priority: $P1 < P2 < \dots < Pn$

For fixed-priority scheduling, the WCRT depends on the alignment of release times




Critical instant

Under FP, for any task τ_i , **the longest response time** occurs when its release coincide with the release of **all higher-priority tasks**.



Critical instant

For **independent preemptive** tasks under fixed priorities, a **critical instant** of τ_i occurs when it releases a job together with all higher-priority tasks and subsequent jobs are released as fast as possible and execute for their WCET.




$$P_1 > P_2 > P_3 > \dots > P_i$$

Notations: $\tau_i(C_i, T_i, D_i, \phi_i)$

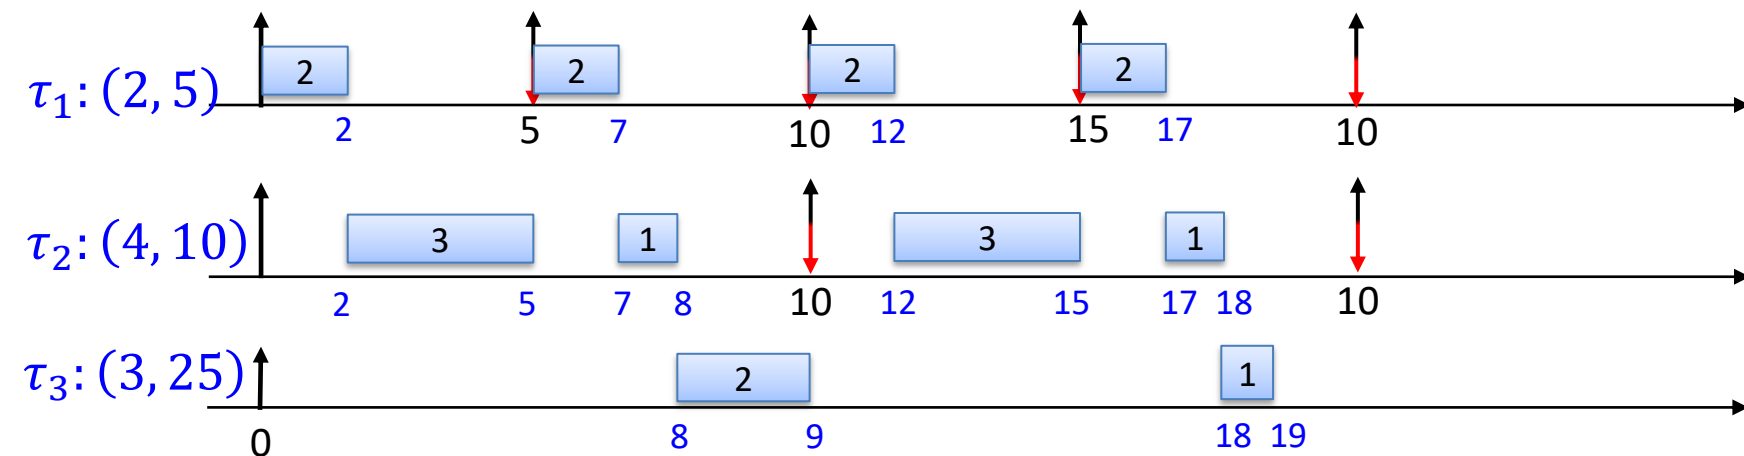
Building an equation for the WCRT

Step 1: If we know that $R_i = X$, then what is the maximum number of jobs of a high-priority task τ_k ($k < i$) that can interfere with τ_i ?



The maximum number of jobs of task τ_k released in an interval of length X is $\left\lceil \frac{X}{T_k} \right\rceil$



Ceiling operator: $\lceil x \rceil = \min\{m \in \mathbb{Z} \mid m \geq x\}$



Building an equation for the WCRT


Step 2: What is the maximum workload that can be generated by a **higher-priority task** τ_k in an interval with length X ?

The maximum workload of task τ_k in X is $\left\lceil \frac{X}{T_k} \right\rceil \cdot C_k$



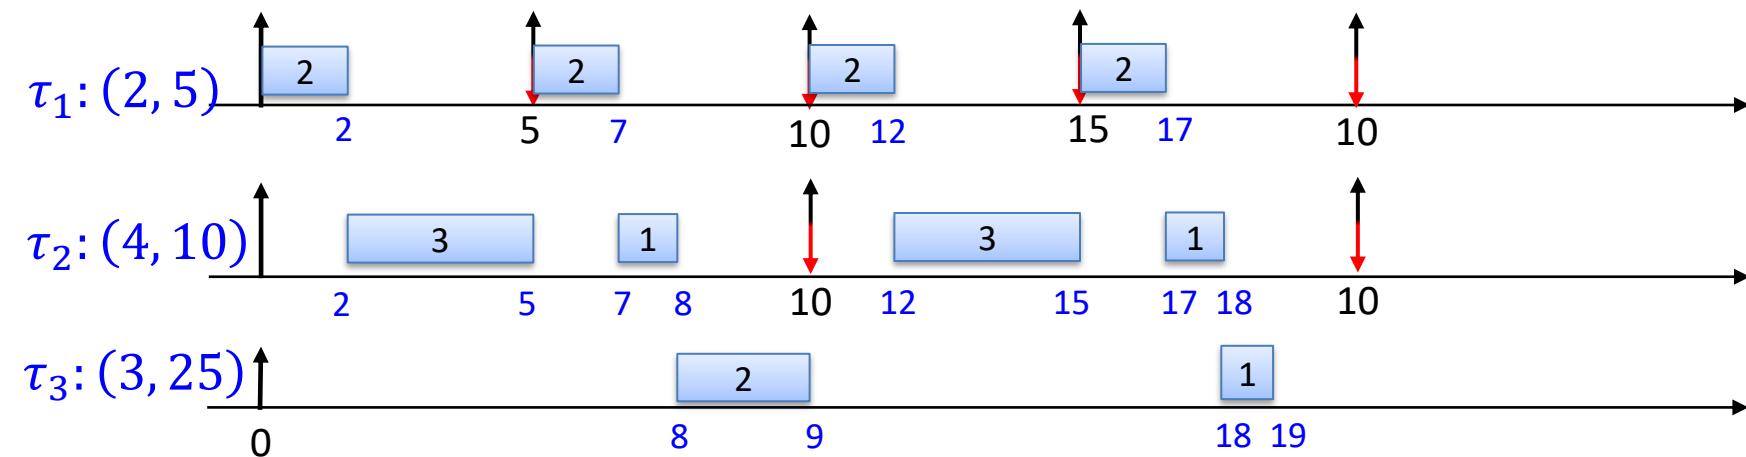
Building an equation for the WCRT

Step 3: What is the maximum workload generated by ALL higher-priority tasks than τ_i in an interval of length X ?




The maximum workload of higher-priority tasks than τ_i is

$$\sum_{k=1}^{i-1} \left\lceil \frac{X}{T_k} \right\rceil \cdot C_k$$



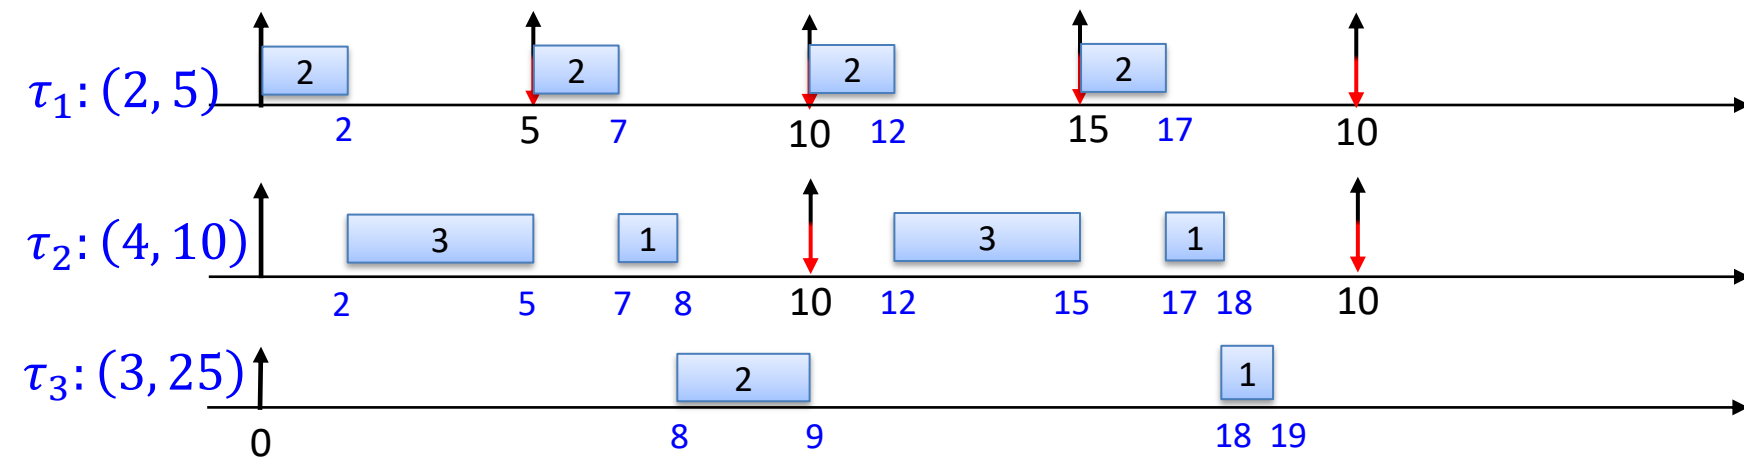
Building an equation for the WCRT

Step 4: What is the maximum workload that must be executed by higher-priority tasks and τ_i itself in an interval of length X ?



The maximum workload of higher-priority tasks and τ_i is

$$C_i + \sum_{k=1}^{i-1} \left\lceil \frac{X}{T_k} \right\rceil \cdot C_k$$



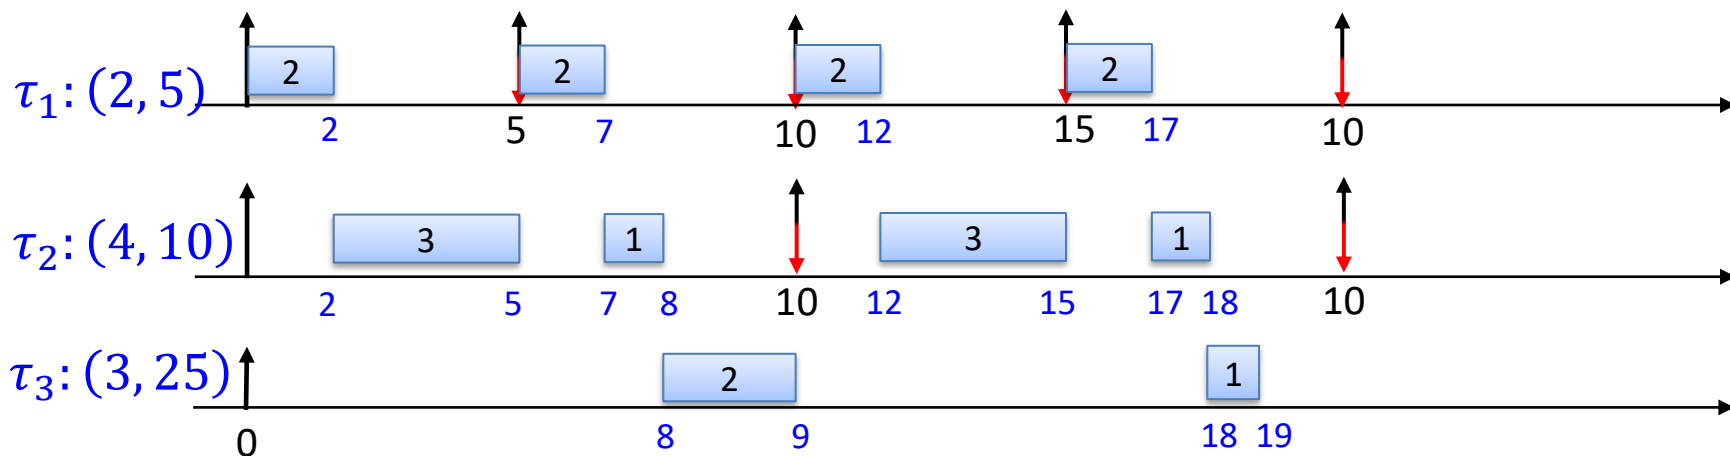

Building an equation for the WCRT

Step 5: How do we find $X = R_i$?

What we execute in X must be equal to the length of $X \rightarrow$

$$X = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{X}{T_k} \right\rceil \cdot C_k$$

Answer: we cannot find it directly!
We need to use a fixed-point iteration method!



On the search for X

1. The response time is at least as large as the WCET, so $X \geq C_i$.

$$R_i^{(0)} = C_i$$

2. Figure out how much workload must be executed in X

$$R_i^{(n)} \leftarrow C_i + \sum_{k=1}^{i-1} \left\lceil \frac{R_i^{(n-1)}}{T_k} \right\rceil \cdot C_k$$

$$C_i + \sum_{k=1}^{i-1} \left\lceil \frac{X}{T_k} \right\rceil \cdot C_k$$

3. Is the workload MORE than the length of X ?
 1. Yes: the WCRT of τ_i must have been larger than X
 2. No: Great! So X is a safe upper bound on the WCRT of τ_i

If $R_i^{(n)} > R_i^{(n-1)}$ then continue and find $R_i^{(n+1)}$
Otherwise stop: $R_i^{(n)}$ is the WCRT

Response-time analysis (RTA)

(Audsley '90)

$$R_i = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{R_i}{T_k} \right\rceil \cdot C_k$$

The solution is based on **fixed-point iterations**:

$$R_i^{(n)} = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{R_i^{(n-1)}}{T_k} \right\rceil \cdot C_k$$

Starting point:

$$R_i^{(0)} = C_i$$

Iterate until:

$$R_i^{(n)} \leq R_i^{(n-1)}$$

Usage: Use $R_i^{(0)}$ to calculate $R_i^{(1)}$, then use $R_i^{(1)}$ to obtain $R_i^{(2)}$, ..., continue until $R_i^{(n)} = R_i^{(n-1)}$

Understanding the terms

Make sure that the execution of τ_i finishes

For all higher-priority tasks

Multiply that by the WCET of τ_k

Count the maximum number of jobs released by task τ_k in the interval $[0, R_i^{(n-1)})$

$$\left\{ \begin{array}{l} R_i^{(n)} = C_i + \sum_{k=1}^{i-1} \left\lfloor \frac{R_i^{(n-1)}}{T_k} \right\rfloor \cdot C_k \\ R_i^{(0)} = C_i \\ \text{Continue if } R_i^{(n)} > R_i^{(n-1)} \end{array} \right.$$


Agenda

- RM utilization-based tests (reminder)
- Response-time analysis for periodic or sporadic tasks under FP
 - **Response-time analysis**
 - For $\forall \tau_i, \sigma_i = 0$ and $R_i \leq T_i$
 - **For $\forall \tau_i, \sigma_i \geq 0$ and $R_i \leq T_i$**
 - For $\forall \tau_i, \sigma_i \geq 0$ and $R_i \leq T_i$ or $R_i > T_i$
 - **Park test ($\forall \tau_i, \sigma_i = 0$ and $R_i \leq T_i$)**

Maximizing interference

What is the **maximum number of jobs** that a **periodic or sporadic** task τ_i with **release jitter** may release in an **interval of duration L** ?

At most $\left\lceil \frac{L + \sigma_i}{T_i} \right\rceil$ jobs



Maximizing interference


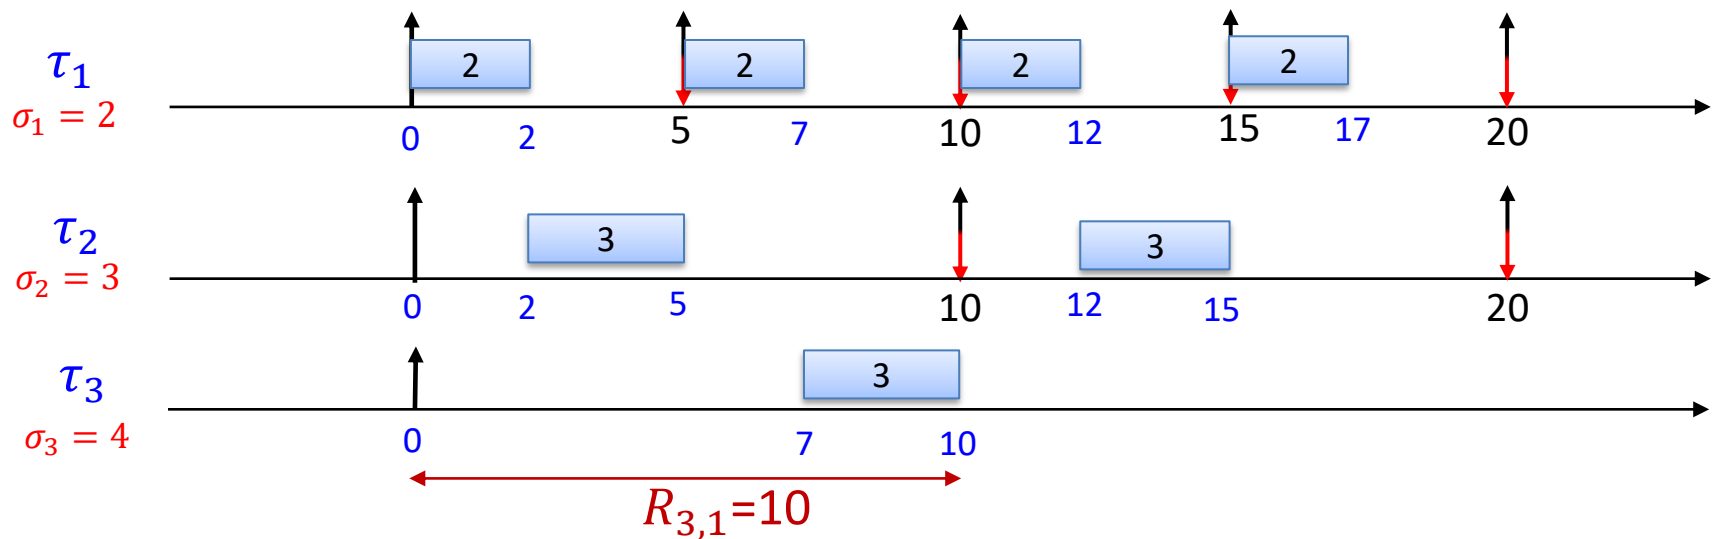
What is the **maximum number of jobs** that a **periodic or sporadic** task τ_i with **release jitter** may release in an **interval of duration L** ?

At most $\left\lceil \frac{L + \sigma_i}{T_i} \right\rceil$ jobs


What is the **maximum workload** that a periodic or sporadic task τ_i with **release jitter** may release in an interval of duration L ?

At most $\left\lceil \frac{L + \sigma_i}{T_i} \right\rceil \times C_i$



Visualizing the critical instant



Visualizing the critical instant




Visualizing the critical instant



Visualizing the critical instant

For **independent preemptive** tasks under fixed priorities, the critical instant of τ_i occurs when it releases a job together with all higher-priority tasks. The first job of each task is **released with maximum jitter**, subsequent jobs are released as fast as possible and execute for their WCET.



Response-time analysis

$$R_i = \sigma_i + X_i$$

where

$$X_i = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{X_i + \sigma_k}{T_k} \right\rceil \cdot C_k$$

The solution is based on **fixed-point iterations**:

$$X_i^{(n)} = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{X_i^{(n-1)} + \sigma_k}{T_k} \right\rceil \cdot C_k$$

Starting point:

$$X_i^{(0)} = C_i$$

Iterate until:


$$X_i^{(n)} \leq X_i^{(n-1)}$$

Why do we add σ_i to X_i ?

The **interference** of higher priority tasks (captured in X_i) happens only **after τ_i is released**.

The **response time** of τ_i is the **difference between the finish time and arrival time** of τ_i , which happens up to σ_i time units before the release of τ_i .


Visualizing the critical instant




Agenda

- RM utilization-based tests (reminder)
- Response-time analysis for periodic or sporadic tasks under FP
 - **Response-time analysis**
 - For $\forall \tau_i, \sigma_i = 0$ and $R_i \leq T_i$
 - For $\forall \tau_i, \sigma_i \geq 0$ and $R_i \leq T_i$
 - For $\forall \tau_i, \sigma_i \geq 0$ and $R_i \leq T_i$ or $R_i > T_i$
 - Park test ($\forall \tau_i, \sigma_i = 0$ and $R_i \leq T_i$)

Effect of $R_i > T_i$



Effect of $R_i > T_i$



First job of τ_3 generates interference on second job of τ_3


→ The worst-case response time may not be experienced by the first job of τ_i anymore!

Solution: analyze **all jobs** released in a **level- i busy window** starting at a critical instant

Level i busy-window

A **level i busy-window** is a window of time during which the processor is **busy** executing tasks with a **priority** equal to or higher than that of τ_i .

To find the WCRT of τ_i , we must calculate the response time of all jobs in the level- i busy window



Level i busy-window

A **level i busy-window** is a window of time during which the processor is **busy** executing tasks with a **priority** equal to or higher than that of τ_i .

The **maximum length of level- i busy window** is given by the first positive solution to:

$$L_i = \sum_{k=1}^i \left\lceil \frac{L_i + \sigma_k}{T_k} \right\rceil \cdot C_k$$

The solution is based on **fixed-point iterations**:

$$L_i^{(n)} = \sum_{k=1}^i \left\lceil \frac{L_i^{(n-1)} + \sigma_k}{T_k} \right\rceil \cdot C_k$$

Starting point:

$$L_i^{(0)} = \sum_{k=1}^i C_k$$

Iterate until:

$$L_i^{(n)} \leq L_i^{(n-1)}$$

Level i busy-window

A **level i busy-window** is a window of time during which the processor is **busy** executing tasks with a **priority** equal to or higher than that of τ_i .

The **maximum length of level- i busy window** is given by the first positive solution to:

$$L_i = \sum_{k=1}^i \left\lceil \frac{L_i + \sigma_k}{T_k} \right\rceil \cdot C_k$$

What is the **maximum** number of jobs released by τ_i in the longest level- i busy window?

$$N_i = \left\lceil \frac{L_i + \sigma_i}{T_i} \right\rceil$$

RTA when $R_i > T_i$

Goal:

Calculate the response time of every job of τ_i in the level-i busy window

1. Calculate the **arrival time** of the j^{th} job of τ_i in the level-i busy window (assuming the first job arrive at time 0):

$$a_{i,j} = (j - 1) \times T_i$$

2. Calculate the worst-case **finish time** of the j^{th} job of τ_i in the level-i busy window:

$$f_{i,j} = X_{i,j} + \sigma_i$$

where

$$X_{i,j} = (j \times C_i) + \sum_{k=1}^{i-1} \left\lceil \frac{X_{i,j} + \sigma_k}{T_k} \right\rceil \cdot C_k$$

3. Calculate the worst-case **response time** of the j^{th} job of τ_i :

$$R_{i,j} = f_{i,j} - a_{i,j}$$

4. Calculate the worst-case **response time** of τ_i :

$$R_i = \max_{\forall k, 1 \leq j \leq N_i} (R_{i,j})$$

Understanding the terms

$$X_{i,j} = j \times C_i + \sum_{k=1}^{i-1} \left\lfloor \frac{X_{i,j} + \sigma_k}{T_k} \right\rfloor \cdot C_k$$

j jobs of τ_i must execute until the finish time of the *j*th jobs of τ_i

For all higher-priority tasks

Multiply that by the WCET of τ_k

Time **from the release** of the first job τ_i to the finish time the *j*th jobs of τ_i

Count the maximum number of jobs released by task τ_k in the interval $[0, X_i)$

$$f_{i,j} = X_{i,j} + \sigma_i$$

Time **from the arrival** of the first job τ_i to the finish time the *j*th jobs of τ_i

Agenda

- RM utilization-based tests (reminder)
- Response-time analysis for FP
 - Response-time analysis
 - **Park test ($\forall \tau_i, \sigma_i = 0$ and $R_i \leq T_i$)**

Using RTA to build fast **sufficient** tests when

$$\forall \tau_i, \sigma_i = 0 \text{ and } R_i \leq T_i$$


Park's test

$$\begin{cases} R_i^{(n)} = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{R_i^{(n-1)}}{T_k} \right\rceil \cdot C_k \\ R_i^{(0)} = C_i \\ \text{Continue if } R_i^{(n)} > R_i^{(n-1)} \end{cases}$$

Idea:

Instead of searching for the exact WCRT, just check if the **workload that must be completed before the deadline** is smaller than the deadline of the task

How?



Using RTA to build fast **sufficient** tests when

$$\forall \tau_i, \sigma_i = 0 \text{ and } R_i \leq T_i$$

Park's test

$$\begin{cases} R_i^{(n)} = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{R_i^{(n-1)}}{T_k} \right\rceil \cdot C_k \\ R_i^{(0)} = C_i \\ \text{Continue if } R_i^{(n)} > R_i^{(n-1)} \end{cases}$$

Idea:

Instead of searching for the exact WCRT, just check if the **workload that must be completed before the deadline** is smaller than the deadline of the task

workload that must be completed
before the deadline


$$C_i + \sum_{k=1}^{i-1} \left\lceil \frac{D_i}{T_k} \right\rceil \cdot C_k \leq D_i$$

What is the computational complexity per task?

$O(n)$

Why Park's test is just a sufficient test?

Park's test is only just a **sufficient** test



Is this task set schedulable by RM?

Yes! The critical instant is schedulable.

=> The WCRT of each task is smaller than its deadline

What is the WCRT of τ_3 and τ_4 ?

8 and 9, respectively.

What does Park's test say about τ_4 ?

It says τ_4 will miss its deadline

Why did it happen?

When a higher-priority task is released after the WCRT of the task under study, Park's test includes that task within the workload that must be finished by the deadline! Namely, it upper-approximate the interfering workload!

$$C_4 + \sum_{k=1}^{4-1} \left\lceil \frac{D_4}{T_k} \right\rceil \cdot C_k \leq D_4 \Rightarrow 1 + \left(\left\lceil \frac{10}{5} \right\rceil \cdot 2 \right) + \left(\left\lceil \frac{10}{9} \right\rceil \cdot 3 \right) + \left(\left\lceil \frac{10}{10} \right\rceil \cdot 1 \right) \\ = 1 + 4 + 6 + 1 = 12 \not\leq 10$$

Summary

Schedulability analyses for preemptive FP scheduling:

- Utilization-based tests

- LL-bound: $U \leq U_{lb}$

- Hyperbolic bound: $f(U_1, \dots, U_n) \leq 2$

- Response time analysis $\forall i, R_i \leq D_i$

- Simplified version: Park's test (over-approximates the interfering workload)

- Simplified test for harmonic task tests: $U \leq 1$